

# CS 111 week 1

Tianxiang Li

# About myself

- Email: [tianxiang@cs.ucla.edu](mailto:tianxiang@cs.ucla.edu)
- Office hours: 4-6Pm PDT, Monday, Zoom
  - Feel free to email me if I didn't see you, or if you don't see me in the room
  - FIFO
- Discussion Section 1B, Friday 10AM-12PM PDT
  - Discussions will be recorded and uploaded to CCLE
  - Slides will be uploaded to CCLE
- Questions
  - Pizza posts
  - Project grading questions, please send a private post

# What will be on the discussion section

- Mainly focus on projects
  - Clarify the projects
  - Discuss background knowledge for projects
  - Introduce APIs used for the projects and show code examples
- Q&A
  - Feel free to speak up when you have question, or type in chatbox
  - Feel free to answer any questions
  - Feel free to correct me when there's a mistake

# Project logistics

- Code will be graded on `Inxsrv09.seas.ucla.edu`
  - Create a Seasnet account if you don't have one
    - Username 8 characters, password 8 characters: at least one number, one lowercase, and one uppercase (no special characters)
  - Forgot your account?
    - `help@seas.ucla.edu`
  - `ssh user_name@Inxsrv09.seas.ucla.edu`
- Make sure the gcc version is correct ( `/usr/local/cs/gcc-9.3.0/bin/` )
  - Command to check version:  
*which gcc*
  - If version is incorrect:  
`export PATH= /usr/local/cs/gcc-9.3.0/bin/:$PATH`  
Add the line `PATH=/usr/local/cs/gcc-9.3.0/bin:$PATH` in `~/.profile` or `~/.bash_profile`
  - Safe way to compile  
`/usr/local/cs/gcc-9.3.0/bin /gcc -o lab0 lab0.c`

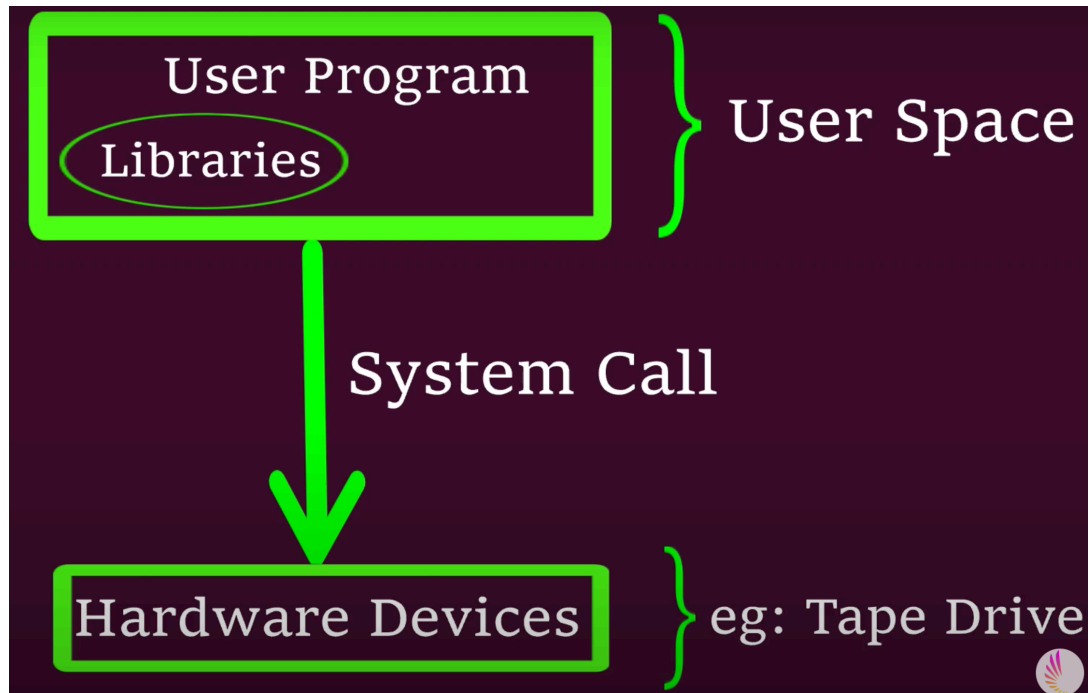
# Project 0: Warm Up

# Project Requirements

- First thing: Write a program (lab0.c)
  - **Copies** its standard **input** to its standard **output** by:
    - **read**(2)-ing from file descriptor 0 (until encountering an end of file) and,
    - **write**(2)-ing to file descriptor 1. If no errors (other than EOF) are encountered,
    - your program should **exit**(2) with a return code of 0.
  - Support **optional** command line **arguments**
    - **--input = *filename*** use the specified file as standard input
    - **--output = *filename*** create the specified file and use it as standard output
    - **--segfault** force segmentation fault
    - **--catch** catch segmentation fault, log error, exit

# System call for file operation: quick review

- Functions provided by Linux to create an interface with the OS
  - `open(2)`, `read(2)`, `write(2)`, `close(2)`
  - `#include <unistd.h>`



# System call for file operation: quick review

- open(2), read(2), write(2), close(2)
- int open(const char \**path*, int *oflags*);
  - e.g. fd = open("/var/tmp/a.out", O\_RDONLY)

ReadOnly: O\_RDONLY  
WriteOnly: O\_WRONLY  
ReadWrite: O\_RDWR

- ssize\_t read(int *fd*, void \**buf*, size\_t *count*);
- ssize\_t write(int *fd*, void \**buf*, size\_t *count*);

*fd*: file descriptor created by open (read from it/write to it)

*buf*: pointer to the buffer (store read results/write using its content)

*count*: size of the buffer

**return value**: Number of bytes read/written, -1 if error occurred

- int close(int fd);  
close fd.



## Code example: copy one file to the other

```
#include <fcntl.h>
```

```
// error handling code omitted for briefly
```

```
#define BUFFERSIZE 1024
```

```
char buffer[BUFFERSIZE];
```

```
int main(argc, char * argv[])
```

```
cp /var/tmp/a.txt /home/user/file.txt
```

```
{
```

```
    src_fd = open(argv[1], O_RDONLY);
```

```
    des_fd = open(argv[2], O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR | S_IRGRP);
```

```
    while ((ret = read(src_fd, buffer, BUFFERSIZE)) > 0)
```

```
        write(des_fd, buffer, ret); //FIXME: not handling actual bytes written < ret.
```

```
    close(src_fd);
```

```
    close(des_fd);
```

```
    return 0;
```

```
}
```

# File Descriptor Review

- File Descriptor: Abstract representation of a file
- 0 → reading from keyboard
- 1 → writing to screen
- 2 → error → screen

# Project overview

- Write a program that **reads** from **stdin** and **writes** to **stdout**, essentially a simplified cat command
- Special file descriptors: 0: stdin, 1: stdout, 2: stderr

- Pseudo-code:

```
while (1)
{
    read(0, buffer);
    if error or reach EOF
        break;
    write(1, buffer);
    if error or reach EOF
        break;
}
```

- File Descriptor: Abstract representation of a file
- 0 → reading from keyboard
- 1 → writing to screen
- 2 → error → screen

# IO re-direction

- Support two options:
  - --input=filename ... use the specified file as standard input
  - --output=filename ... create the specified file and use it as standard output
- One possible way to do it:

```
    if (input option is set)
```

```
        in_fd = open(input_filename);
```

```
    else
```

```
        in_fd = 0;
```

```
    while (1)
```

```
    {
```

```
        read(in_fd, buffer);
```

```
        write(stdout, buffer);
```

```
    }
```

# IO re-direction (cont)

- Support two options:
  - --input=filename ... use the specified file as standard input
  - --output=filename ... create the specified file and use it as standard output
- The other way of doing it: dup(2) system calls:

prototype: `int dup(int fd);`

Make the lowest available file descriptor to point to the same file as fd.


- Read redirection example:

```
int ifd = open("/var/tmp/input.txt", O_RDONLY);
if (ifd >= 0) {
    close(0);
    dup(ifd);
    close(ifd);
}
```

# IO-redirection with dup(2)

- Prototype: `int dup(int fd);`  
Make the lowest available file descriptor to point to the same file as `fd`.
- Read redirection example:


```
int ifd = open("/var/tmp/input.txt", O_RDONLY);
if (ifd >= 0) {
    close(0);
    dup(ifd);
    close(ifd);
}
```



# IO-redirection with dup(2)

- Prototype: `int dup(int fd);`  
Make the lowest available file descriptor to point to the same file as `fd`.
- Read redirection example:


```
int ifd = open("/var/tmp/input.txt", O_RDONLY);
if (ifd >= 0) {
    close(0);
    dup(ifd);
    close(ifd);
}
```



# IO-redirection with dup(2)

- Prototype: `int dup(int fd);`  
Make the lowest available file descriptor to point to the same file as `fd`.
- Read redirection example:

```
int ifd = open("/var/tmp/input.txt", O_RDONLY);
if (ifd >= 0) {
    close(0);
    dup(ifd);
    close(ifd);
}
```






# IO-redirection with dup(2)

- Prototype: `int dup(int fd);`  
Make the lowest available file descriptor to point to the same file as `fd`.
- Read redirection example:

```
int ifd = open("/var/tmp/input.txt", O_RDONLY);  
if (ifd >= 0) {  
    close(0);  
    dup(ifd);  
    close(ifd);  
}
```



# SIGSEGV signal

- Support two options:
  - --segfault: force a segmentation fault
  - --catch: use signal(2) to register a SIGSEGV handler that catches the segmentation fault
- What truly happens upon a segfault:

```
int main(void)
```

```
{
```

```
    char * ptr = NULL;
```

```
    (*ptr) = 0;
```

```
}
```

***NULL pointer dereference***

*when a program attempts to read or write to memory with a NULL pointer.*

- cause an illegal memory access exception
- CPU traps to the illegal memory access exception handler in the OS
- OS sends a SIGSEGV signal to the process
- return from OS to process, process **execute the default SIGSEGV handler** that kills the process (via exit)

# signal(2)

- Support two options:
  - **--segfault**: force a segmentation fault
  - **--catch**: use signal(2) to register a SIGSEGV handler that catches the segmentation fault
- signal(2): register a signal handler to replace the default signal handler

```
#include <signal.h>
```

```
typedef void (*sighandler_t)(int);
```

```
sighandler_t signal(int signum, sighandler_t handler);
```

**signal()** sets the disposition of the signal *signum* to *handler*, which  
SIGINT is a signal generated when a user presses Control-C. This will terminate the program from the terminal.  
SIGSTOP tells LINUX to pause a process to be resumed later.  
SIGSEGV is sent to a process when it has a [segmentation fault](#).  
SIGKILL is sent to a process to cause it to terminate at once.

...

## signal(2) example

```
void sigsegv_handler(int sig)
{
    fprintf(...)
    exit(....)
}
int main(void)
{
    char * ptr = NULL;
    signal(SIGSEGV, sigsegv_handler);
    (*ptr) = 0; // Jump to execute code in sigsegv_handler function
}
```

# Error handling: errno and strerror

- When you print out an error message, your message should include enough information for user to understand its cause
- **errno**: a *global integer variable* which is set by system calls and some library functions in the event of an error to indicate what went wrong.
  - *errno is a global variable defined in <errno.h>, the systems sets it whenever a system call error occurs*

ret = read(src\_fd, buffer, BUFFERSIZE) → ret = -1 → an error occurred

errno is set to tell you what specifically went wrong:

errno = 9 (EBADF): src\_fd is an invalid file descriptor

errno = 14 (EFAULT): buffer points to an invalid memory address

# Error handling: strerror and code example

- When you print out an error message, your message should include enough information for user to understand its cause

- `char * strerror(int errno);`

The `strerror()` function returns a pointer to a string that describes the error code passed in the argument `errno`.

```
int infd = open(infile, O_RDONLY);  
if (infd < 0)  
{  
    fprintf(stderr, "%s: %s\n", infile, strerror(errno));  
}
```

# Handle input arguments

Suppose we have a program called test, which supports two arguments:

--verbose #output the debugging logging message

--input file # specify a file as the input

```
./test --input /home/user/file.txt
```

```
./test --verbose
```

```
./test --input /home/user/file.txt --verbose
```

```
./test --input /home/user/file.txt
```

```
./test --verbose --input #error
```

...

Imagine:

(1) parse command line arguments for programs with tens/hundreds of options.

(2) parse command line arguments for every program you wrote

→ Standard APIs: (e.g. getopt\_long) to handle command line arguments

```

struct option args[] = {
/*      option      has_arg      flag      return value */
{"input",          1,          NULL,      'i'},      /* input file      */
{"verbose",        0,          NULL,      'v'},      /* verbose          */
{ 0, 0, 0, 0 }
};

```

*The getopt\_long() function works like getopt() except that it also accepts long options, started with two dashes.*

*e.g. --arg=param or --arg param.*

```

int main(int argc, char * argv[]) {
    int verbose = 0;
    char* infile;
    while( (i = getopt_long(argc, argv, "", args, NULL) ) != -1) {
        switch(i)
        {
            case 'i':
                infile = optarg; //optarg is a predefined variable, points to the argument of the option
                break;
            case 'v':
                verbose = 1;
                break;
            default:
                error_handling
                break;
        }
    }
}

```



# Makefile

create a Makefile that supports the following targets:

**(default):** build the lab0 executable

**check:** runs a quick smoke-test on whether or not the program seems to work

**clean:** delete all files created by the Makefile

**dist:** build the distribution tarball.

# Makefile review

- Makefile consists of multiple rules, for each rule:

- **Target:** `depend0, depend1, depend2...`

[TAB] command1

[TAB] command2

....

if (1) target does not exist or (2) the modification time of **dependencies** is newer than **target**, then execute the commands. Dependencies can be omitted;

**myprog:** `myprog.c`

`gcc -o myprog myprog.c`

`clean:`

`rm -f myprog`

`$make #execute the first target myprog`

`$make clean`

Makefile skeleton for project 0

lab0: lab0.c:

gcc ...

clean:

rm -f ...

dist: lab0.c and some more files

tar ...

check: lab0 and some more files

...

# GDB

## Task0:

run your program (with the **--segfault** argument) under gdb(1)

take the fault

get a stack **backtrace**

take a screen snapshot (to be included with your submission)

## Task1:

run your program (with the **--segfault** argument) under gdb(1)

set a **break-point** at the bad assignment

**run** the program up to the breakpoint

**inspect the pointer** to confirm that it is indeed NULL

take a screen snapshot (to be included with your submission)

# GDB review

1. Compile the program with `-g` option

```
gcc -g -o myprog myprog.c
```

2. `gdb ./myprog`

3. (gdb) run options

e.g. `run --input=/var/tmp/input.txt`

4. Use various gdb commands to help you debug the program

`backtrace (bt).` → a summary of how your program got where it is

`breakpoint (b) line#` → makes your program stop at a certain point

`print(p) variable_name.` → Prints the value of a given expression

# GDB commands

1. (gdb) bt #show the stacktrace

(gdb) bt

#0 0x080486a2 in func () at test.c:128

#1 0x0804879b in main (argc=2, argv=0xbffff3c4) at test.c:72

2. (gdb) b test.c:14 #set a breakpoint at line 14 of test.c

3. (gdb) print var #print the value of variable var

# Usage Statement

Usage statement usually consists of the correct command line usage for the program and includes a list of the correct command-line arguments or options

usage: myprog

usage: myprog arg1 arg2

usage: lab0 [--input=file] (example: *lab0 --input=file1.txt*)