

CS 111 Week 2

Project 1A: Terminal IO and IPC

Discussion 1B

Tianxiang Li

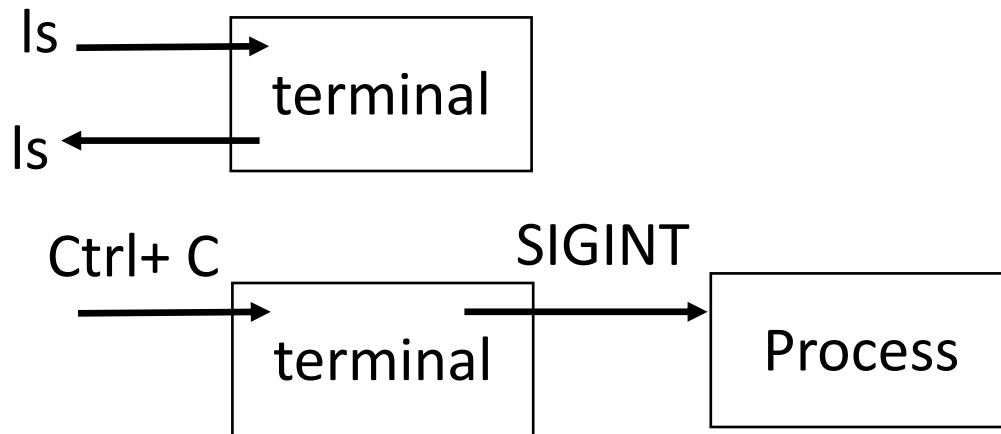
Project 1A

- Terminal
 - Character-at-a-time, full duplex terminal I/O
- Shell
 - Passing input and output between two processes
 - Shutdown processing

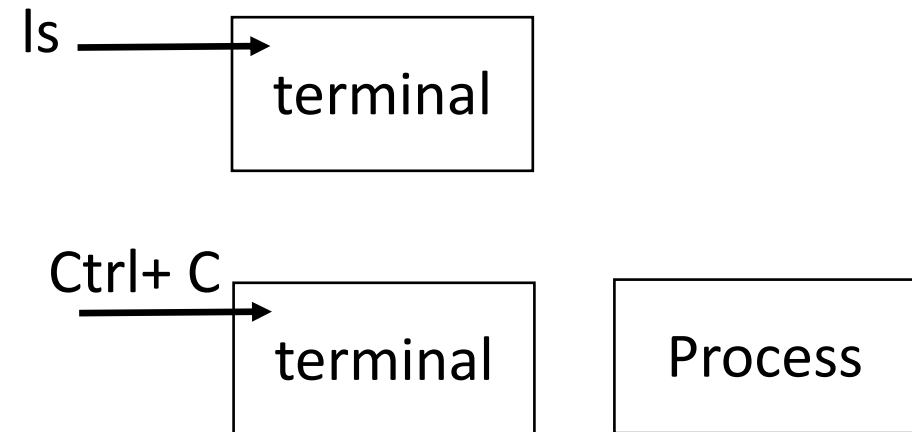
Part A: Terminal

1. Set the stdin into no echo mode, non-canonical input
2. Read from stdin, write to stdout
3. If read '\r' or '\n', write '\r' and '\n' to stdout.
4. When receive EOF (ctrl + D), restore the terminal to normal mode and exit

default terminal mode:
echo and canonical



no echo and non-canonical mode:
implement echo and canonical mode with
our own code




Non-canonical input, no echo mode: APIs

struct termios


```
{  
    tcflag_t c_iflag;    /* input modes */  
    tcflag_t c_oflag;    /* output modes */  
    tcflag_t c_cflag;    /* control modes */  
    tcflag_t c_lflag;    /* local modes */  
    cc_t      c_cc[NCCS]; /* special characters */  
}
```

`int tcgetattr(int fd, struct termios *termios_p);` // **get the state** of the terminal specified by fd



stores state in termios structure

`int tcsetattr(int fd, int optional_actions, const struct termios *termios_p);`
// **set the state** of the terminal specified by fd
// optional_actions: how to set it (e.g. waits for all buffered input/output to finish?)



set params from termios structure

Returns 0 if success and -1 if non-success

Non-canonical input, no echo mode: example

```
void terminal_setup(void) //error handling code omitted for brevity
```

```
{
```

```
    struct termios tmp;
```

```
    tcgetattr(0, &tmp);
```

```
    tmp.c_iflag = ISTRIP, tmp.c_oflag = 0, tmp.c_lflag = 0;
```

```
    tcsetattr(0, TCSANOW, &tmp);
```

```
}
```

change occurs immediately on start-up (before any characters are entered)

```
//set character-at-a-time no-echo mode
```

```
    //tcgetattr → get terminal mode, save normal mode
```

```
    //struct termios → specify required mode
```

```
    //tcsetattr → set mode
```

```
//restore saved terminal mode and exit
```

```
    //tcsetattr → restore mode
```

```
    //exit
```

Part A: Terminal

1. Set the stdin into non-canonical input, no echo mode
2. Read from stdin, write to stdout
3. If read '\r' or '\n', write '\r' and '\n' to stdout.
4. If read EOF (ctrl + D), restore the terminal to normal mode and exit

Part A: Terminal skeleton code

```
int main(void)
{
    terminal_setup();
    while (1)
    {
        read(0, buffer, sizeof(buffer));
        for (char in buffer)
        {
            if (char == 0x4) //ctrl + D
                write(1, "^D"), restore terminal, exit.
            else if (char == '\r' || char == '\n')
                write(1, "\r\n");
            else
                write(1, char);
        }
    }
}
```

Part B: Shell

Extend your program to support a new command line option: `--shell=program` (*--shell=/bin/bash*)

1. Fork a new process to execute program

2. Forward the stdin of the original process to the stdin of the new process.

'\r' or '\n' should be forwarded as '\n'

3. Forward the stdout/stderr of the new process to the stdout of the original process.

'\n' should be forwarded as '\r'

1. When receive ctrl + C, send a SIGINT signal to the new process

2. When

- A. receive EOF from the stdin of the original process

- B. receive EOF from the stdout of the original process

- C. SIGPIPE signal

Get all the output from the new process, report the exit status of the new process and exit.

Fork APIs

`pid_t fork(void)`

creates a new process (child process) by duplicating the calling process (parent process).

Child process has (almost) the same state as the parent process, including register/memory state and file descriptor.

On success, return 0 to the child process

return pid of the child process to the parent process

Both parent and child process start executing from the next instruction

```
int main(void)
{
    fork();
    printf("Hello World\n");
}
```

Fork example

```
int main(void)
```

```
{
```

```
    int ret = fork();
```

```
    if (ret == 0)
```

```
    {
```

```
        printf("Child\n");
```

```
    }
```

```
    else if (ret > 0)
```

```
    {
```

```
        printf("Parent\n");
```

```
    }
```

```
}
```

//If ret < 0 means error occurred, e.g. system out of memory

Parent process

Printf("Parent\n")

CPU0

Child process

Printf("Child\n")

CPU1

Concurrent execution!
non-deterministic results!

What is the output of this program?

waitpid APIs

`pid_t waitpid(pid_t pid, int *wstatus, int options);`

Block and wait for the termination of the **child process**.

Exit status of the child process will be stored at `wstatus`.

Low order 7-bits of `wstatus`: signal number → `WTERMSIG(wstatus)`

The next higher order byte of `wstatus`: exit code → `WEXITSTATUS(wstatus)`

```
SHELL EXIT SIGNAL=0 STATUS=0
```

- After you have closed the write pipe to the shell and processed the final output returned from the shell:
 - use `waitpid(2)` to await the process' completion and capture its return status

Fork and waitpid example

```
int main(void)
{
    int ret = fork();
    if (ret == 0)
    {
        printf("Child process\n");
        exit(5);
    }
    else if (ret > 0)
    {
        int status = 0;
        waitpid(ret, &status, 0);
        printf("Child process exits with code: %d\n", WEXITSTATUS(status));
    }
}
```

Fork and waitpid example

```
int main(void) {  
    int ret = fork();  
    if (ret == 0) {  
        printf("child!\n");  
        exit(5);  
    }  
    else if (ret > 0) {  
        int status = 0;  
        waitpid(ret, &status, 0);  
        printf("Child process exits  
        with: %d\n", WEXITSTATUS(status));  
    }  
}
```

What is the output of this program?

Parent process

waitpid: blocked

waitpid: return

printf(...);

Parent process

waitpid: wakeup

printf(...);

Child process

printf("child!\n");

exit(5);

Child process

printf("child!\n");

exit(5);

waitpid: non-deterministic
→ deterministic

APIs to create a new process

```
int execlp(const char *file, const char *arg, ... /* (char *) NULL */);
```

Replaces the current process image with a **new** process image.

file: path to the new process image (e.g. your executable file name)

arg: command line arguments of the new process, NULL terminated.

arg[0] is the same as file name.

Example:

```
execlp( shell_pgm, shell_pgm, NULL );
```

Fork, exec and waitpid example

```
int main(void)
```

```
{
```

```
    int ret = fork();
```

```
    if (ret == 0)
```

```
    {
```

```
        execlp("ls", "ls", "-a", "-l", NULL);
```

```
        exit(5);
```

```
    }
```

```
    else if (ret > 0)
```

```
    {
```

```
        int status = 0;
```

```
        waitpid(ret, &status, 0);
```

```
        printf("Child process exits with code: %d\n", (status & 0xff));
```

```
    }
```

```
}
```

Parent

```
waitpid(...)  
printf(...)
```

Child

```
execlp(..)  
exit(5)
```

What's the value of
WEXITSTATUS(status)?

Fork, exec and waitpid example

```
int main(void)
```

```
{
```

```
    int ret = fork();
```

```
    if (ret == 0)
```

```
    {
```

```
        execlp("ls", "ls", "-a", "-l", NULL);
```

```
        exit(5);
```

```
    }
```

```
    else if (ret > 0)
```

```
    {
```

```
        int status = 0;
```

```
        waitpid(ret, &status, 0);
```

```
        printf("Child process exits with code: %d\n", (status & 0xff));
```

```
    }
```

```
}
```

Parent

```
waitpid(...)  
printf(...)
```

Child

```
ls -a -l
```

What's the value of
WEXITSTATUS(status)?

0 because execlp starts a new
program

Part B: Shell

Extend your program to support a new command line option: `--shell=program`

1. Fork a new process to execute program
2. Forward the stdin of the original process to the stdin of the new process. `'\r'` or `'\n'` should be forwarded as `'\n'`
3. Forward the stdout/stderr of the new process to the stdout of the original process. `'\n'` should be forwarded as `'\r'`
4. When receive ctrl + C, send a SIGINT signal to the new process
5. When
 - A. receive EOF from the stdin of the original process
 - B. receive EOF/Error from the stdout/stderr of the new process
 - C. receive SIGPIPE signal

Get all the output from the new process, report the exit status of the new process and exit.

Pipe syscalls

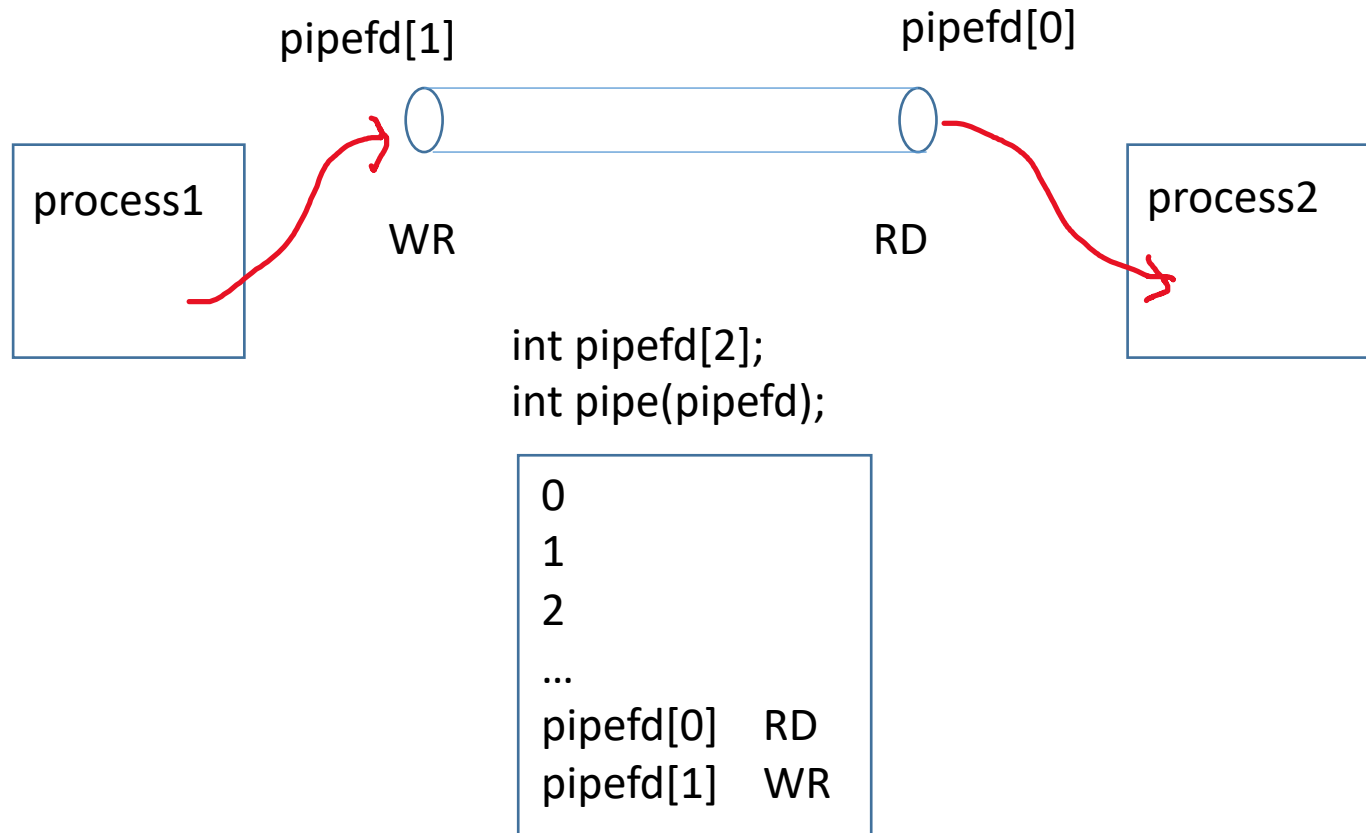
```
int pipefd[2];
```

```
int pipe(pipefd);
```

Creates a pipe, a *unidirectional* data channel (*inter-process communication*).

pipefd returns two file descriptor.

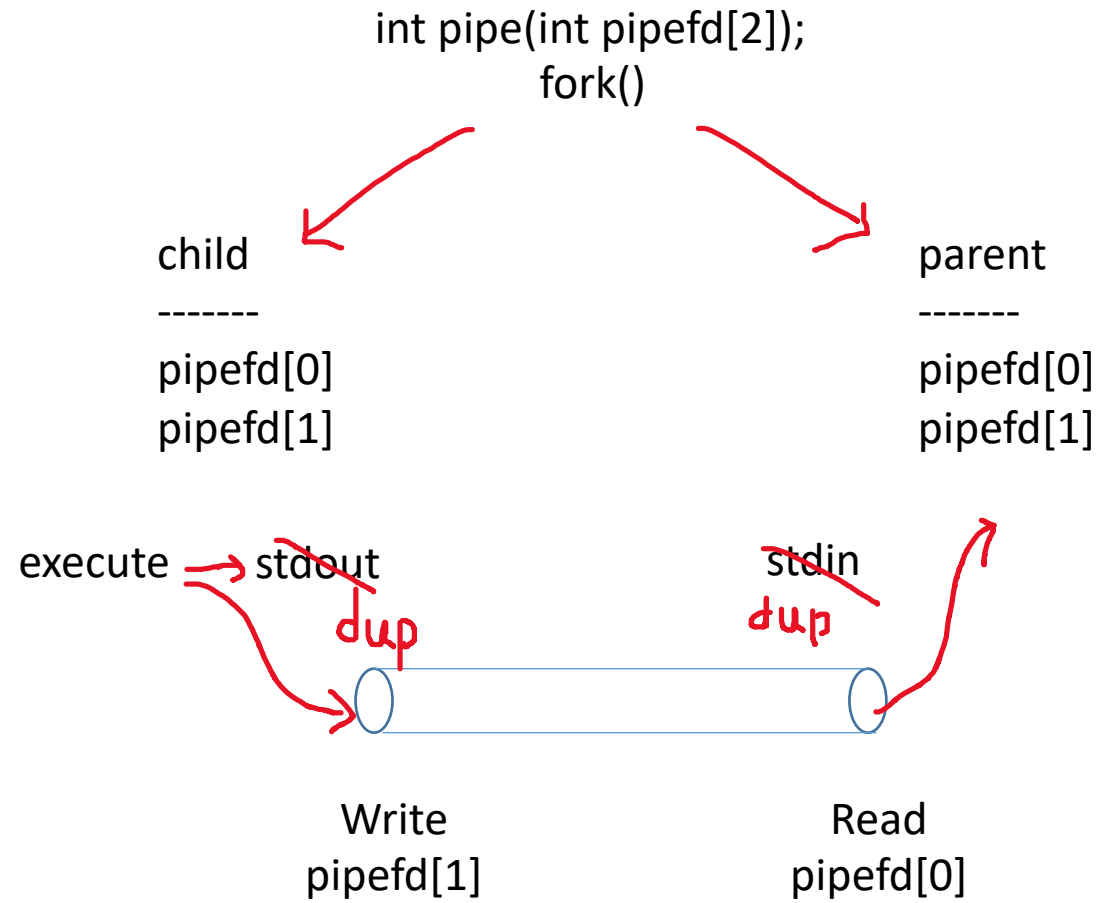
pipefd[0] is the **read** end of the pipe and **pipefd[1]** is the **write** end of the pipe



Pipe syscall examples

```
int main(void)
{
    int fds[2];
    pipe(fds); //pass the pointer of fds array to pipe syscall
               // once returns: fds[0] for read end of pipe, fds[1] for write end of pipe
    write(fds[1], "hahaha", 6);
    read(fds[0], buffer , 6); //buffer will be "hahaha"
}
```

pipe and fork

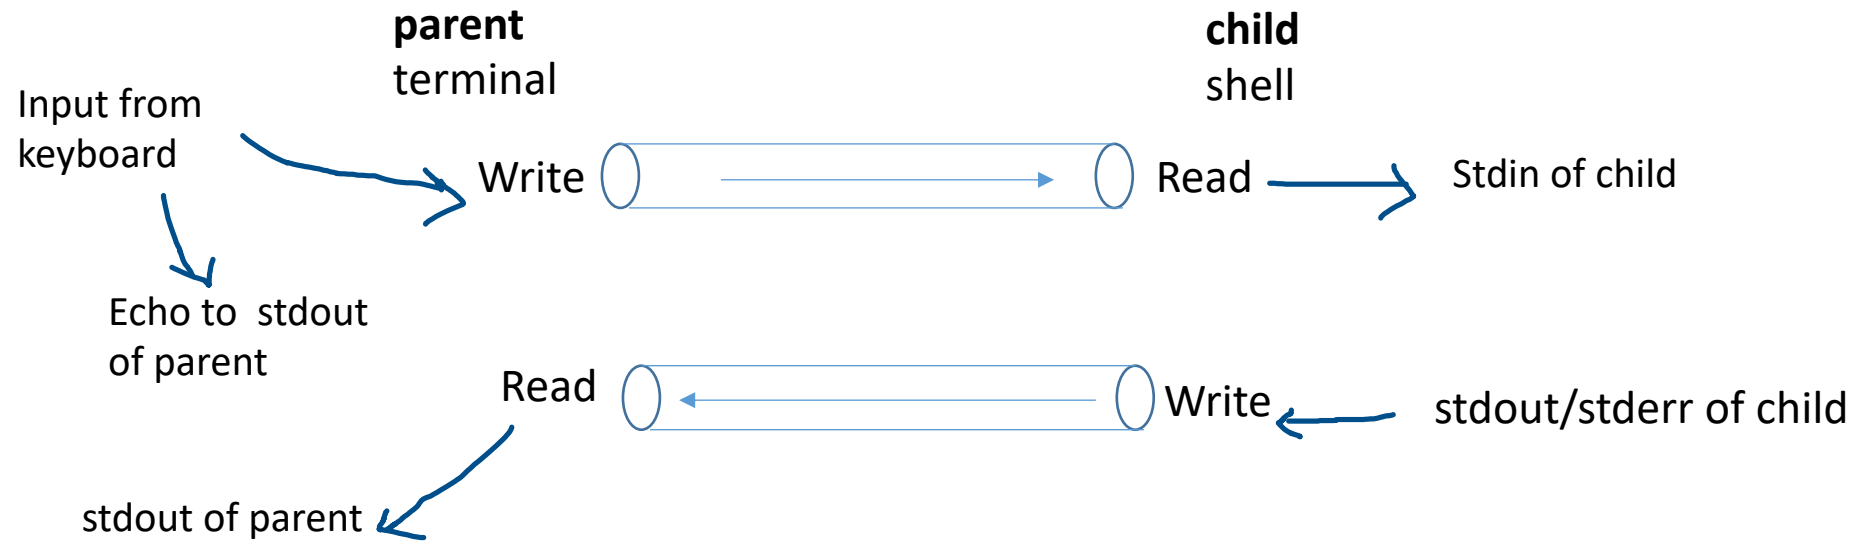


Pipe syscall examples

```
int main(void)
{
    int fds[2];
    pipe(fds);
    ret = fork();
    if (ret == 0)
    {
        read(fds[0], buffer, 6); //buffer is "hahaha"
    }
    else if (ret > 0)
    {
        write(fds[1], "hahaha", 6);
    }
}
```

Forward stdin/stdout between two process

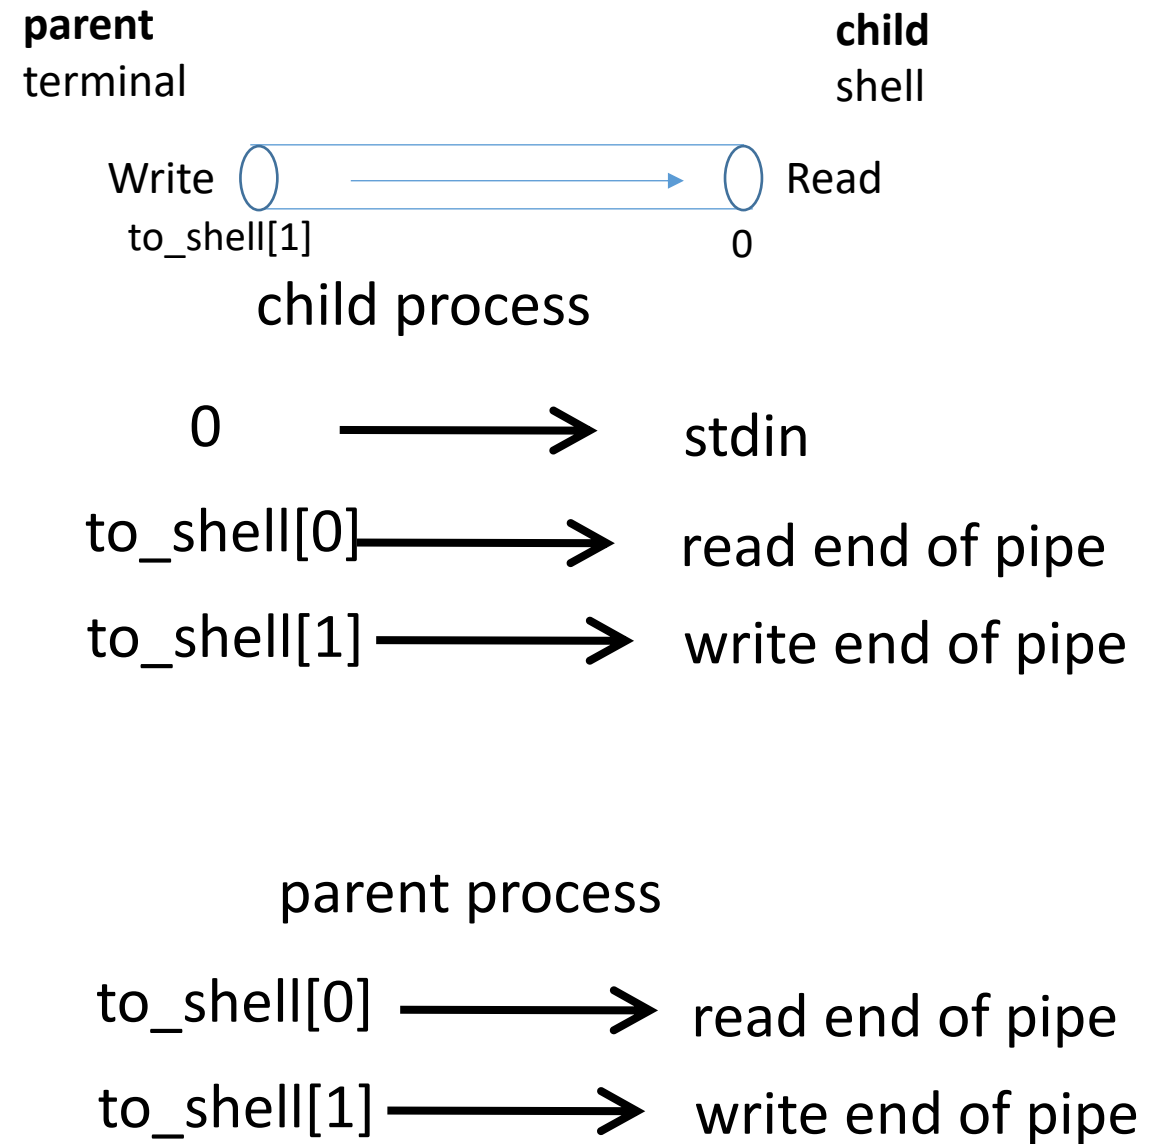
Create two pipes, one end is the stdin, stdout/stderr of the child process, the parent process writes/reads from the other end of the pipe.



Forward stdin/stdout between two process

Create two pipes, one end is the stdin, stdout/stderr of the child process, the parent process writes/reads from the other end of the pipe.

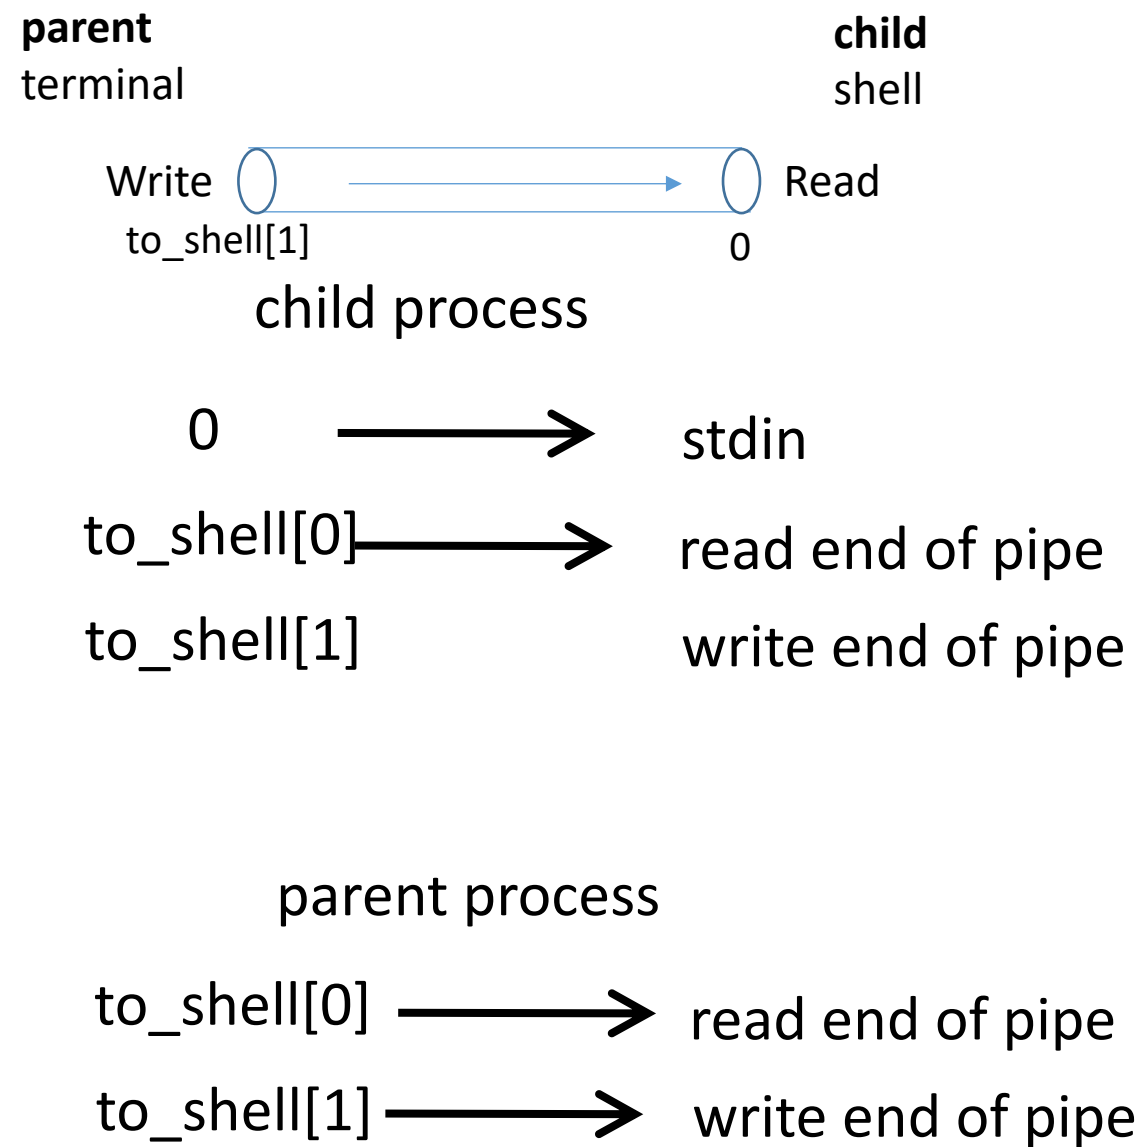
```
int to_shell[2];
pipe(to_shell);
if (fork() == 0) { //child process
    close(to_shell[1]);
    close(0);
    dup(to_shell[0]); //0 → read end of the pipe
    close(to_shell[0]);
    execlp(...);
}
else { //parent process
    close(to_shell[0]);
    while (read(0, buffer, sizeof(buffer)) > 0)
        write(to_shell[1], buffer, sizeof(buffer));
}
```



Forward stdin/stdout between two process

Create two pipes, one end is the stdin, stdout/stderr of the child process, the parent process writes/reads from the other end of the pipe.

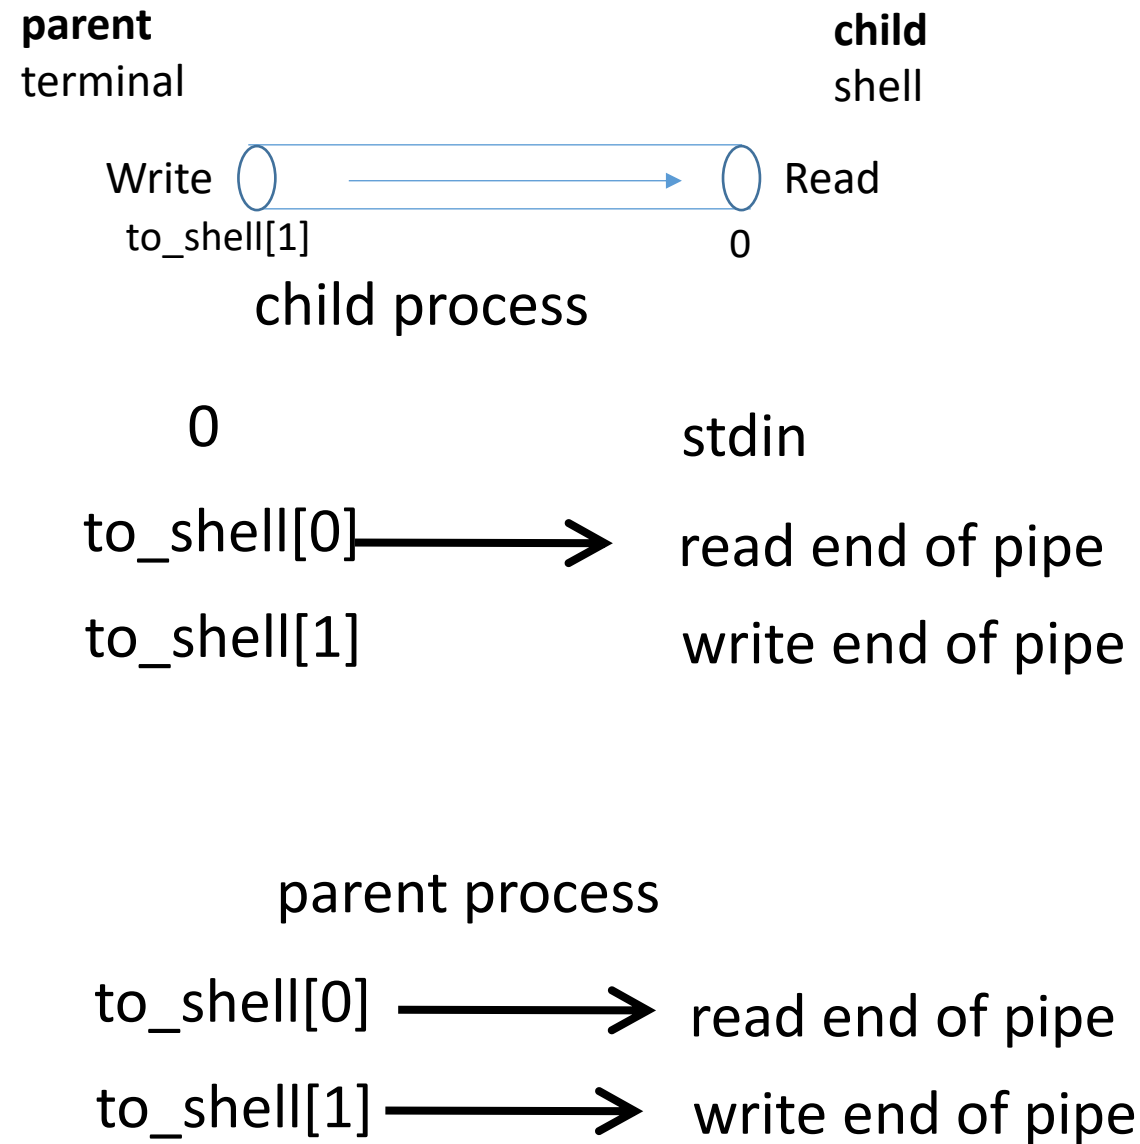
```
int to_shell[2];
pipe(to_shell);
if (fork() == 0) { //child process
    close(to_shell[1]);
    close(0);
    dup(to_shell[0]); //0 → read end of the pipe
    close(to_shell[0]);
    execlp(...);
}
else { //parent process
    close(to_shell[0]);
    while (read(0, buffer, sizeof(buffer)) > 0)
        write(to_shell[1], buffer, sizeof(buffer));
}
```



Forward stdin/stdout between two process

Create two pipes, one end is the stdin, stdout/stderr of the child process, the parent process writes/reads from the other end of the pipe.

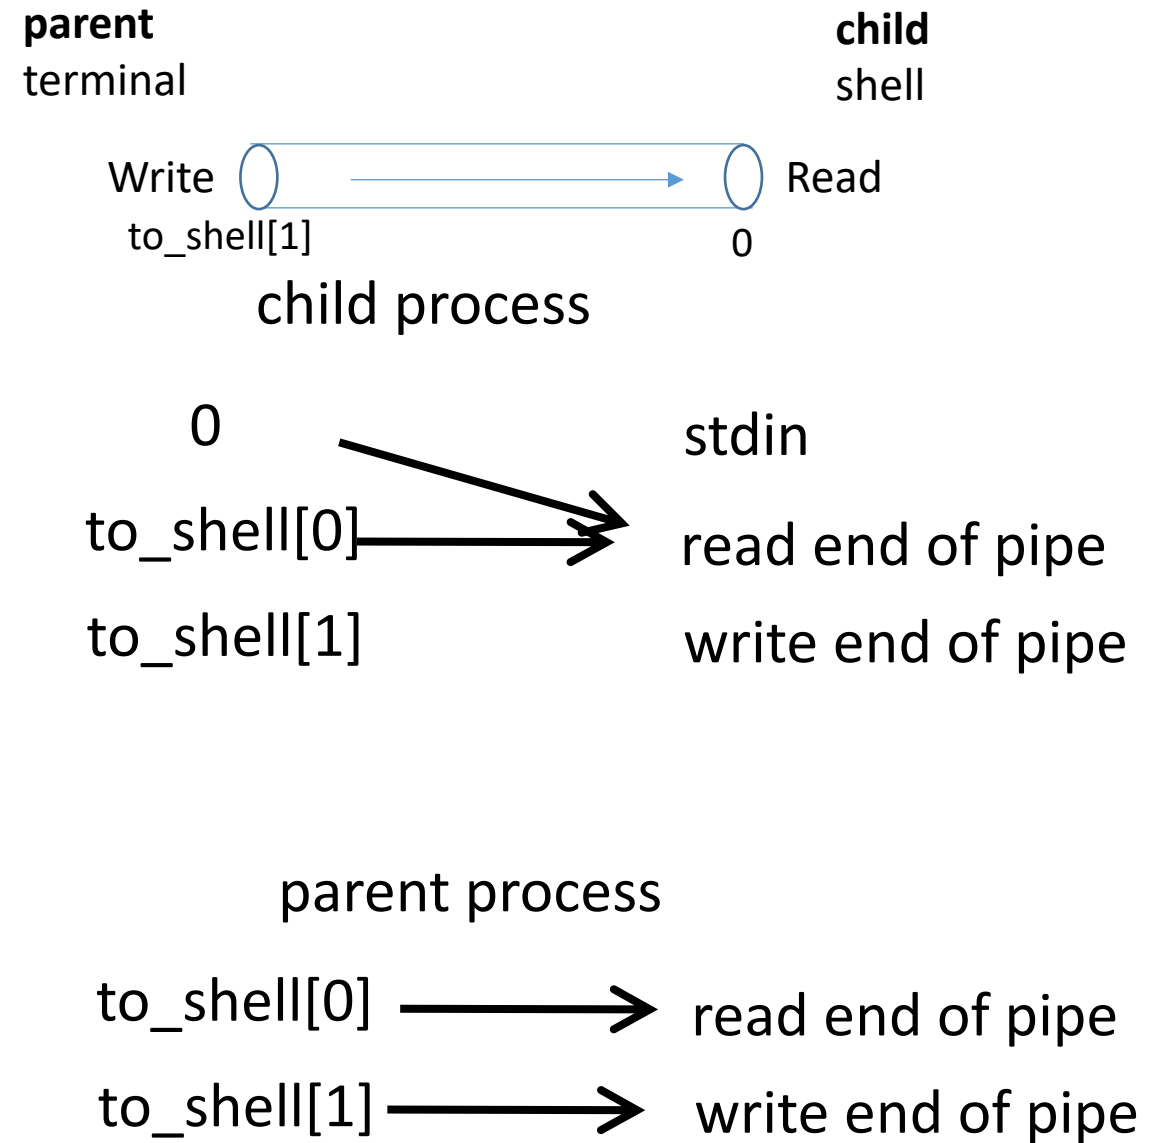
```
int to_shell[2];
pipe(to_shell);
if (fork() == 0) { //child process
    close(to_shell[1]);
    close(0);
    dup(to_shell[0]); //0 → read end of the pipe
    close(to_shell[0]);
    execlp(...);
}
else { //parent process
    close(to_shell[0]);
    while (read(0, buffer, sizeof(buffer)) > 0)
        write(to_shell[1], buffer, sizeof(buffer));
}
```



Forward stdin/stdout between two process

Create two pipes, one end is the stdin, stdout/stderr of the child process, the parent process writes/reads from the other end of the pipe.

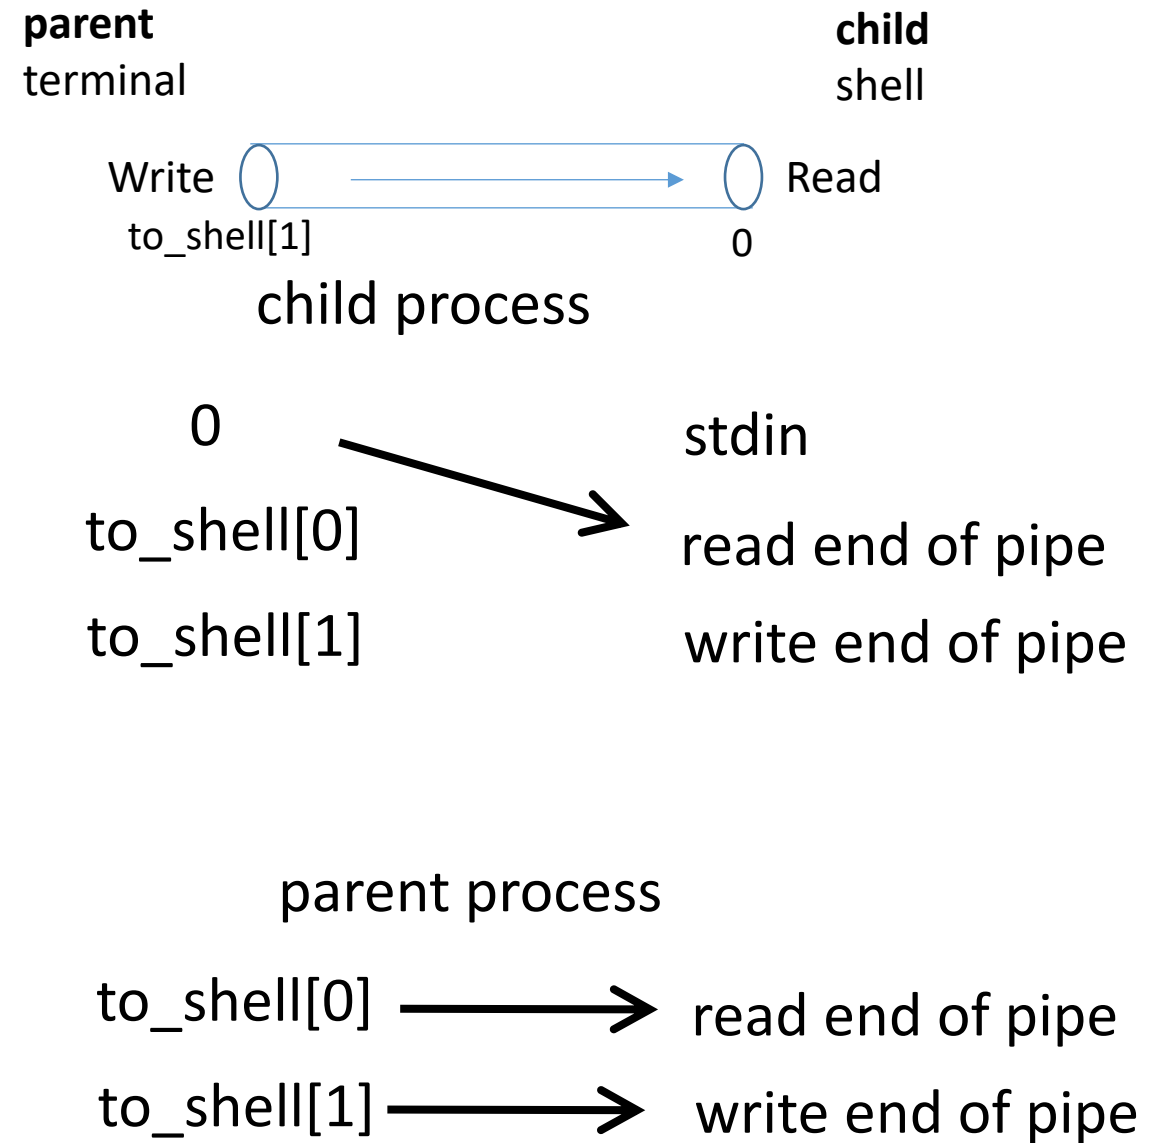
```
int to_shell[2];
pipe(to_shell);
if (fork() == 0) { //child process
    close(to_shell[1]);
    close(0);
    dup(to_shell[0]); //0 → read end of the pipe
    close(to_shell[0]);
    execlp(...);
}
else { //parent process
    close(to_shell[0]);
    while (read(0, buffer, sizeof(buffer)) > 0)
        write(to_shell[1], buffer, sizeof(buffer));
}
```



Forward stdin/stdout between two process

Create two pipes, one end is the stdin, stdout/stderr of the child process, the parent process writes/reads from the other end of the pipe.

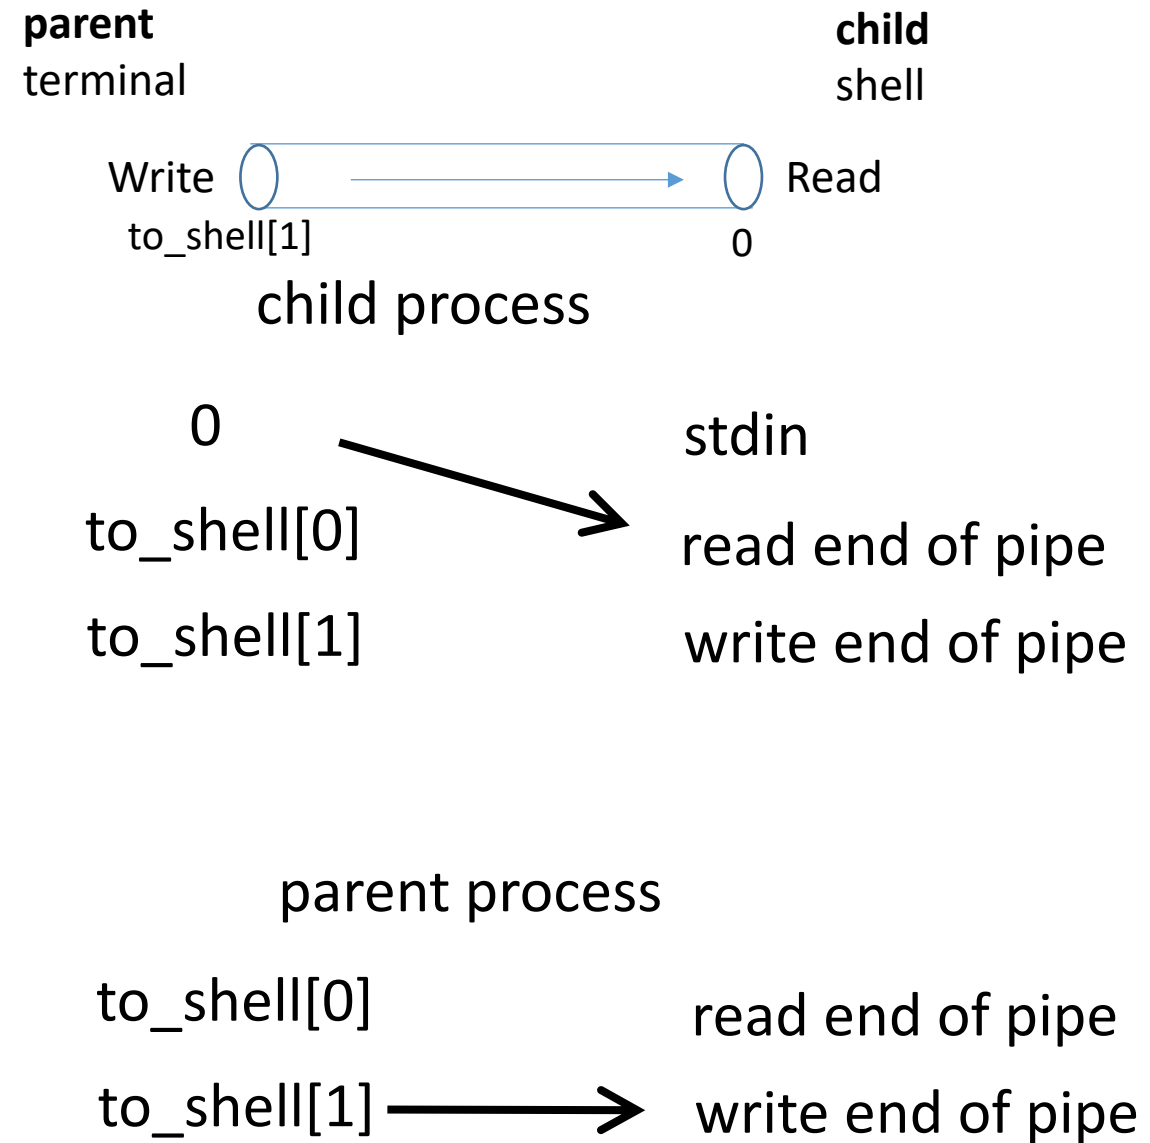
```
int to_shell[2];
pipe(to_shell);
if (fork() == 0) { //child process
    close(to_shell[1]);
    close(0);
    dup(to_shell[0]); //0 → read end of the pipe
    close(to_shell[0]);
    execlp(...);
}
else { //parent process
    close(to_shell[0]);
    while (read(0, buffer, sizeof(buffer)) > 0)
        write(to_shell[1], buffer, sizeof(buffer));
}
```



Forward stdin/stdout between two process

Create two pipes, one end is the stdin, stdout/stderr of the child process, the parent process writes/reads from the other end of the pipe.

```
int to_shell[2];
pipe(to_shell);
if (fork() == 0) { //child process
    close(to_shell[1]);
    close(0);
    dup(to_shell[0]); //0 → read end of the pipe
    close(to_shell[0]);
    execlp(...);
}
else { //parent process
    close(to_shell[0]);
    while (read(0, buffer, sizeof(buffer)) > 0)
        write(to_shell[1], buffer, sizeof(buffer));
}
```



```
handle_cmd_line(argc, argv); // prog = "new program to execute"
```

Create pipes

```
if (fork() == 0) { //child
```

```
    close unused end of pipes
```

```
    redirect stdin, stdout, stderr, to/from pipes.
```

```
    execvp(prog, prog, NULL);
```

```
}
```

```
else { //parent process
```

```
    close unused end of pipes
```

```
    while (1) {
```

```
        read(stdin, buffer, sizeof(buffer)); //input from keyboard
```

```
        handle "\r", "\n", forward to pipe, stdout
```

```
        read(from_shell[0], buffer, sizeof(buffer)); //input from shell
```

```
        handle "\n", forward to stdout
```

```
    }
```

```
}
```

Problem: Read may be
blocked forever!

Solution: Make sure read will
not block before reading it.

poll APIs

```
int poll(struct pollfd *fds, nfds_t nfds, int timeout);
```

poll blocks and returns when

- (1) one of a set of file descriptors to become ready to perform I/O or
- (2) the specified timeout (number of milliseconds) has passed.

fds: an array of struct pollfd

nfds: number of elements in fds array

```
struct pollfd {  
    int  fd;           /* file descriptor to monitor */  
    short events;      /* events to monitor, e.g. ready to read, errors occurred */  
    short revents;     /* return value, encode which events have occurred */  
};
```

```

pollfds[0].fd = 0;
pollfds[0].events = POLLIN + POLLHUP + POLLERR;
pollfds[1].fd = from_shell[0];
pollfds[1].events = POLLIN + POLLHUP + POLLERR;

while (1) {
    poll(pollfds, 2, -1); // -1 means no time out.
    if (pollfds[0].revents & POLLIN) {
        count = read(0, buf, sizeof(buf));
        forward to stdout, to_shell [1]
    }
    if (pollfds[1].revents & POLLIN) {
        count = read(from_shell[0], buf, sizeof(buf));
        forward to stdout.
    }
    if (pollfds[0].revents & (POLLHUP | POLLERR)) {
    }
    ....
}

```

POLLIN: data to read,
POLLHUP: FD is closed by the other side,
POLLERR: error occurred

```

struct pollfd {
    int fd;          /* file descriptor to monitor */
    short events;    /* events to monitor, e.g. ready to read, errors occurred */
    short revents;   /* return value, encode which events have occurred */
};
    
```

Part B: Shell

Extend your program to support a new command line option: `--shell=program`

1. Fork a new process to execute program
2. Forward the stdin of the original process to the stdin of the new process. `'\r'` or `'\n'` should be forwarded as `'\n'`
3. Forward the stdout/stderr of the new process to the stdout of the original process. `'\n'` should be forwarded as `'\r'`
4. When receive ctrl + C, send a SIGINT signal to the new process
5. When
 - A. receive EOF from the stdin of the original process
 - B. receive EOF/Error from the stdout/stderr of the new process
 - C. receive SIGPIPE signal

Get all the output from the new process, report the exit status of the new process and exit.

kill syscalls APIs

```
int kill(pid_t pid, int sig);
```

sends the signal specified by sig to pid.

```
read(0, buffer, sizeof(buffer));
```

```
for (char in buffer)
```

```
{
```

```
    if (char == 0x03) //Recieve ctrl + C
```

```
    {
```

```
        write(1, "^C", 2);
```

```
        kill(child_pid, SIGINT);
```

```
    }
```

```
}
```

Part B: Shell (Error Handling)

When

- A. receive EOF from the stdin of the original process
- B. receive EOF/Error from the stdout/stderr of the new process
- C. receive SIGPIPE signal

Get all the **output** from the new process, report the **exit status** of the new process and **exit**.

For A: Detect ^D character in buffer

For B: Handled by poll: revents & (POLLHUP | POLLERR)

For C: Need to register **SIGPIPE** handler in parent process

indication that the shell has shut down