

CS 111 week 4

Project 1B: Compressed Network Communication

Discussion 1B

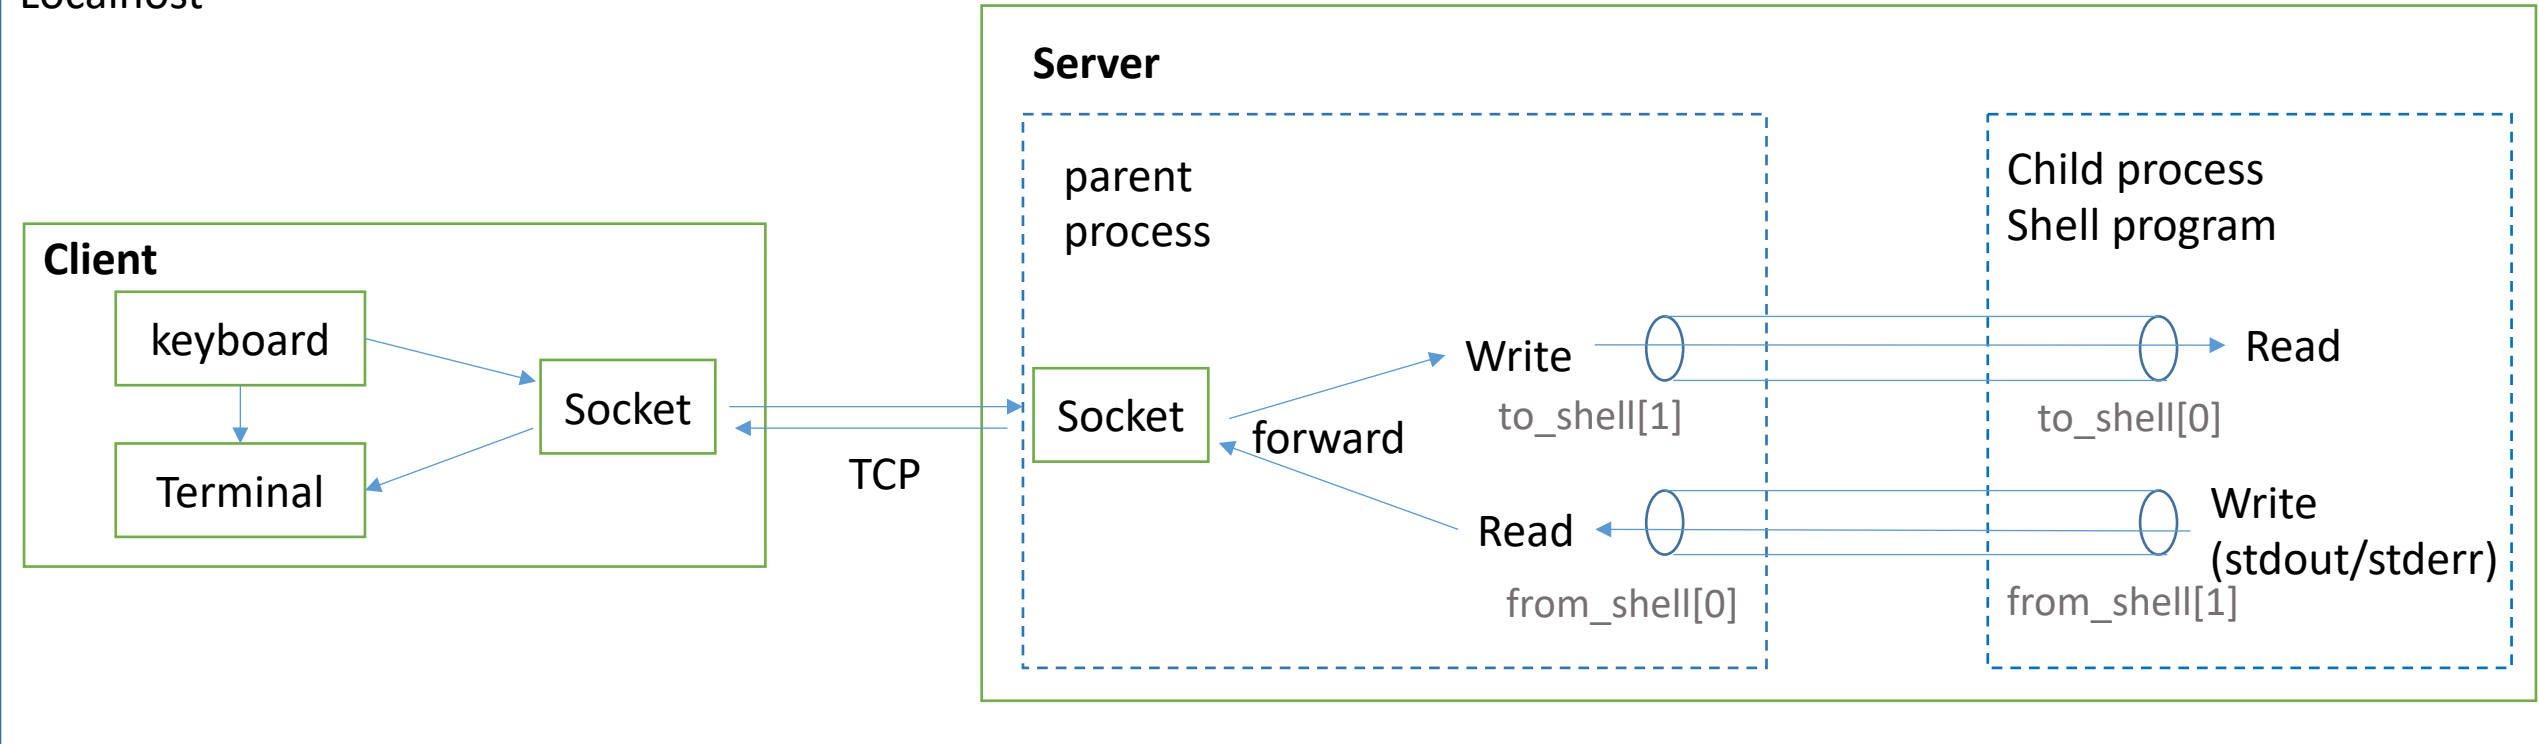
Tianxiang Li

Project Overview


1. Part A: extend the project in Lab1A with network connection.
2. PartB: Compress the network connections

Project Overview

Localhost




Background: C/S Model, IP/Port




- IP address identifies which host the data is destined to
- Port number identifies which services on the host the data is destined to
 - e.g. HTTP 80, SSH 22, FTP DATA 20, MySQL 3306, ... (0-1024 reserved)

Background: Socket



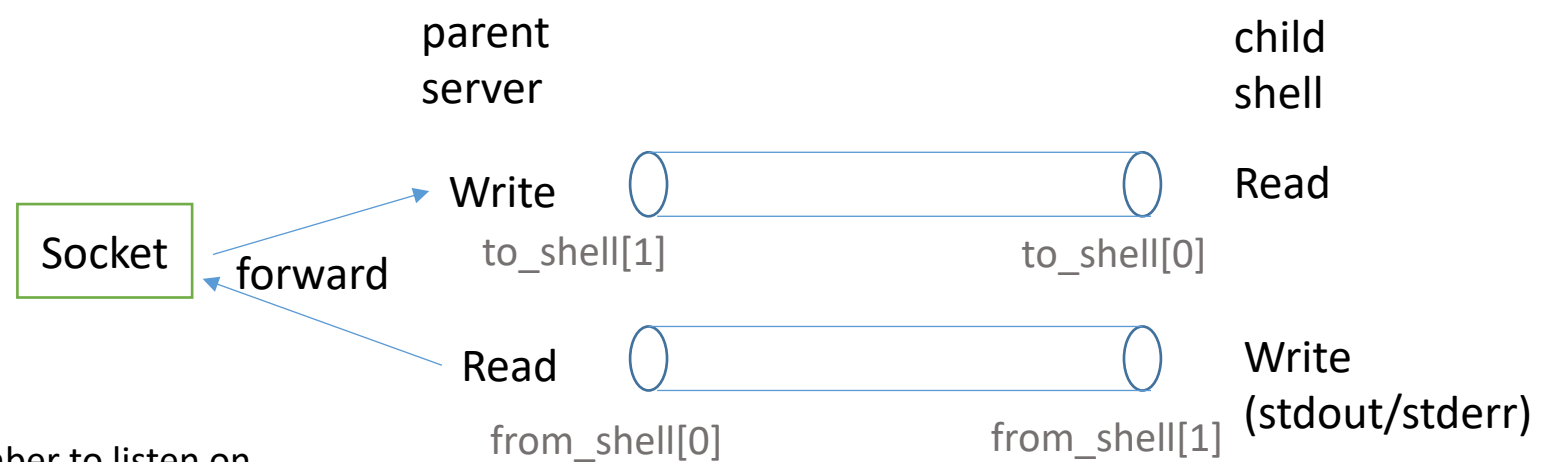
- **A socket is “one endpoint of a two-way communication link between two programs running on the network.”**
 - IP address identifies which host the data is destined to
 - Port number identifies which services on the host the data is destined to
 - e.g. HTTP 80, SSH 22, FTP DATA 20, MySQL 3306, ... (0-1024 reserved)

Client



- Options
 - **--port=portnum** (mandatory): port number used to set up connection
 - **--log=filename**: maintains a record of data sent over the socket
 - **--compress** (second part of project)
- I/O
 - Use **non-canonical** (character at a time, no echo) terminal behavior you used in Project 1A.
 - send input from the **keyboard** to the **socket** (while echoing to the **display**)
 - send input from the **socket** to the **display**
- Special character handling
 - If a **^D** or **^C** is entered on the terminal, simply pass it through to the server like any other character.
- Error handling
 - **unrecognized argument**: print an informative **usage message** (on **stderr**) and exit with a **return code of 1**.
 - **system call fails**: print an informative **error message** (on **stderr**) and exit with a **return code of 1**.
- Exit
 - Before exiting, **restore** normal terminal modes.

Server



- Options
 - **--port=portnum** (mandatory): port number to listen on
 - **--shell** (can be omitted) fork a child process, which will exec a shell to process commands received
 - **--compress** (second part of project)
- I/O
 - Redirect the shell process's stdin/stdout/stderr to the appropriate pipe ends (similar to what you did in Project 1A).
 - Input received through the network **socket** should be forwarded through the **pipe to the shell**.
 - Input received from the **shell pipes** (stdout/stderr from the shell) should be forwarded out to the network **socket**.
- Special character handling
 - Because the server forks the shell (and knows its process ID), the processing of **^C** (turning it into a **SIGINT** to the **shell**) must be done in the server process.
 - when the server encounters a **^D** it should **close** the **write side of the pipe** to the shell (as was done in project 1A).
- Error handling
 - Receive **EOF** or **SIGPIPE** from the **shell pipes** (e.g. the shell exits): harvest the shell's completion **status**, log it to **stderr** (as you did in project 1A), and exit with a **return code of 0**.
 - Receive **EOF** or **read error** from the **network connection**, (e.g., the client exits): **close** the **write pipe** to the shell, **await** an **EOF** from the pipe from the shell, and harvest and report its termination **status** (as above).
 - **unrecognized argument**: print an informative **usage message** (on **stderr**) and exit with a **return code of 1**.
 - **system call fails**: print an informative **error message** (on **stderr**) and exit with a **return code of 1**.
- Exit
 - Before exiting, **restore** normal terminal modes.

Testing

- Start Server
 - Listen on port greater than 1024
 - Do not display anything (other than error messages) on the server side
- Start Client
 - Connect to same port as server
 - You will see your commands echoed back to the terminal as well as the output from the shell

Shutdown

- Server (Server initiates shutdown)
 - an *exit(1)* command is sent to the shell, or the server closes the write pipe to the shell.
 - the **shell exits**, causing the server to receive an **EOF** on the read pipe from the shell.
 - the server collects and reports the shell's **termination status**.
 - the server **closes** the network **socket** to the client, and **exits**.
- Client
 - the client continues to process output from the server until it receives an error on the network socket from the server.
 - the client **restores** terminal modes and **exits**.

TCP socket: service setup

TCP Client

TCP Server



TCP socket: service setup

TCP Client

TCP Server


socket()

socket (): Create a socket

TCP socket: service setup

TCP Client

TCP Server




bind(): bind a socket to a local IP address and port #

TCP socket: service setup

TCP Client

TCP Server




listen(): passively waiting for connections

TCP socket: service setup

TCP Client



TCP Server



accept(): accept a new connection

TCP socket: service setup

TCP Client



TCP Server

socket()



bind()



listen()



accept()



blocked until
connection
from client

TCP socket: service setup

TCP Client

socket()

TCP Server

socket()

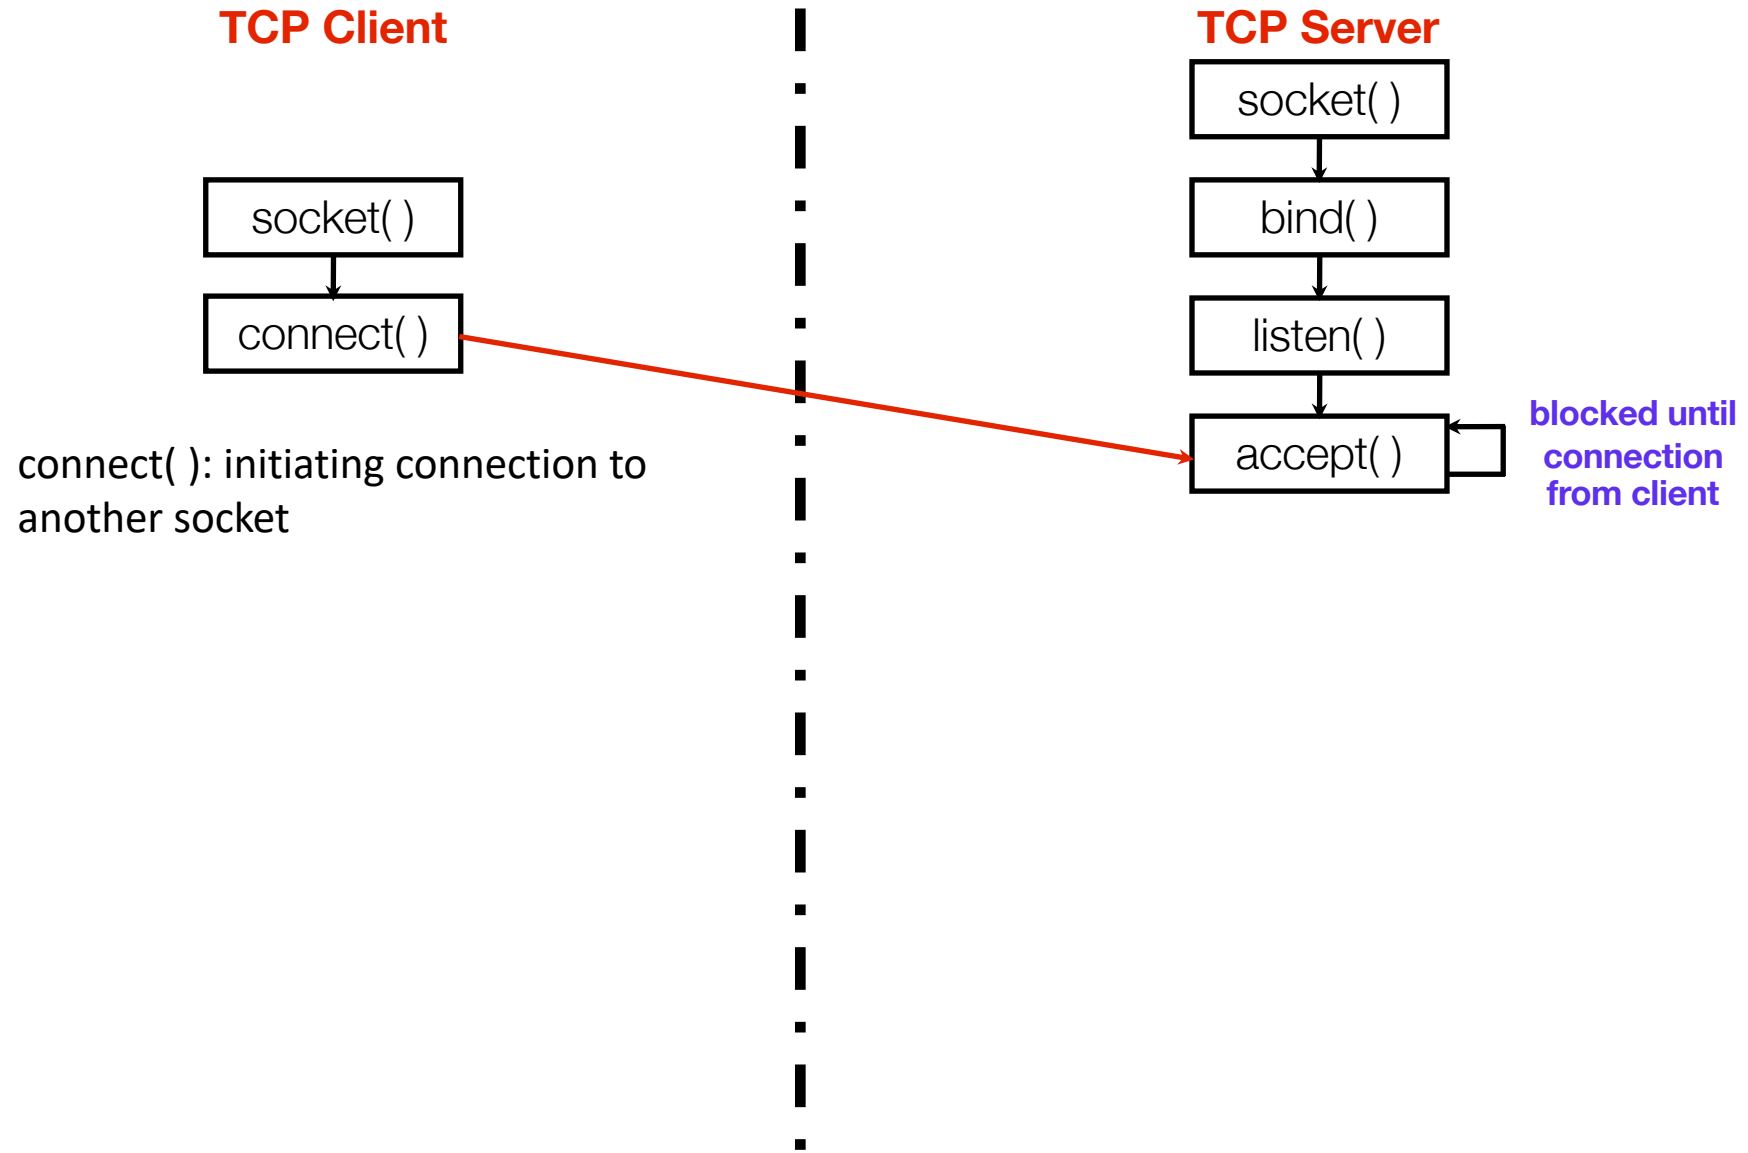
bind()

listen()

accept()


blocked until
connection
from client

TCP socket: establish connection




TCP socket: send and receive data

TCP Client



TCP Server

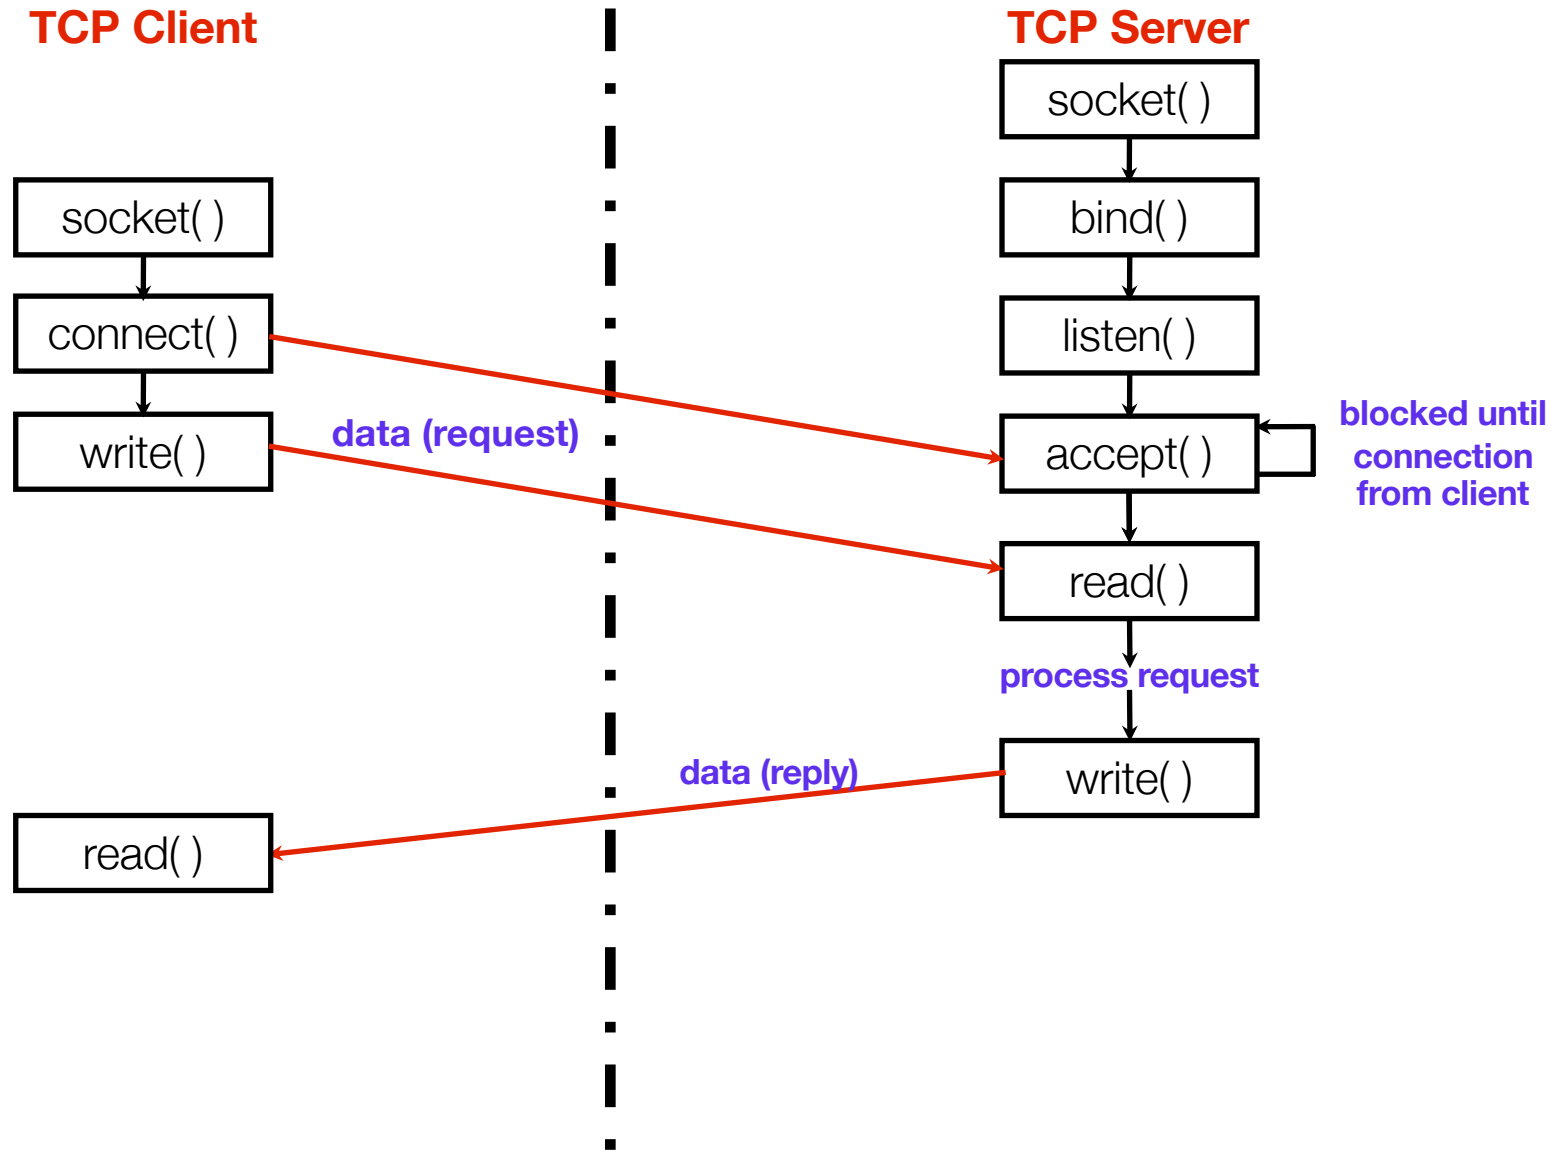


data (request)

blocked until
connection
from client

process request

TCP socket: send and receive data



Headers

```
/* headers */  
#include <stdio.h>  
#include <stdlib.h>  
#include <string.h>  
#include <errno.h>  
#include <sys/types.h>  
#include <sys/socket.h>  
#include <sys/wait.h>  
#include <netinet/in.h>  
...
```

How to write a client?

```
/* include all the headers */  
int client_connect(char* hostname, unsigned int port) {  
  
    //1. create socket  
  
    //2. fill in socket address information  
  
    //3. connect socket with corresponding address  
  
}
```

Functions

- **int socket(int *socket_family*, int *type*, int *protocol*);**
 - Create a socket
 - returns the socket descriptor or -1(failure). Also sets errno on failure
 - ***socket_family***: protocol family
 - **AF_INET** for IPv4, **AF_INET6** for IPv6, ...
 - ***type***: communication style
 - **SOCK_STREAM** for TCP
 - **SOCK_DGRAM** for UDP
 - ***protocol***: protocol within family, which is typically set to 0

How to write a client?

```
/* include all the headers */
int client_connect(char* hostname, unsigned int port) {
/* e.g. host_name:"google.com", port:80, return the socket for subsequent communication */

    int sockfd;
    struct sockaddr_in serv_addr; /* server addr and port info */
    struct hostent* server;
    //1. create socket
    sockfd = socket(AF_INET, SOCK_STREAM, 0);

    //2. fill in socket address information

    //3. connect socket with corresponding address
    return sockfd;
}
```

Socket programming API: essential structs

- sockfd — socket descriptor. Just a regular int.
- sockaddr_in — struct for socket address info

```
struct sockaddr_in { // used for IPv4 only
    short      sin_family; // addr family, AF_INET
    unsigned short sin_port; // port number
    struct in_addr sin_addr; // internet address
    unsigned char sin_zero[8]; // padding zeros
};
```


How to write a client?

```
struct sockaddr_in { // used for IPv4 only
    short    sin_family; // addr family, AF_INET
    unsigned short sin_port; // port number
    struct in_addr sin_addr; // internet address
    unsigned char sin_zero[8]; // padding zeros
};
```

```
/* include all the headers */
int client_connect(char* hostname, unsigned int port) {
    /* e.g. host_name:"google.com", port:80, return the socket for subsequent communication */

    int sockfd;
    struct sockaddr_in serv_addr; /* server addr and port info */
    struct hostent* server;
    //1. create socket
    sockfd = socket(AF_INET, SOCK_STREAM, 0);

    //2. fill in socket address information
    serv_addr.sin_family = AF_INET;
    serv_addr.sin_port = htons(port);
    server = gethostbyname(hostname); /* convert host_name to IP addr */
    memcpy(&serv_addr.sin_addr.s_addr, server->h_addr, server->h_length); /* copy ip address from server to serv_addr */
    memset(serv_addr.sin_zero, '\0', sizeof serv_addr.sin_zero); /* padding zeros */

    //3. connect socket with corresponding address

    return sockfd;
}
```

How to write a client?

```
/* include all the headers */
int client_connect(char* hostname, unsigned int port) {
/* e.g. host_name:"google.com", port:80, return the socket for subsequent communication */

    int sockfd;
    struct sockaddr_in serv_addr; /* server addr and port info */
    struct hostent* server;
    //1. create socket
    sockfd = socket(AF_INET, SOCK_STREAM, 0);

    //2. fill in socket address information
    serv_addr.sin_family = AF_INET;
    serv_addr.sin_port = htons(port);
    server = gethostbyname(hostname); /* convert host_name to IP addr */
    memcpy(&serv_addr.sin_addr.s_addr, server->h_addr, server->h_length); /* copy ip address from server to serv_addr */
    memset(serv_addr.sin_zero, '\0', sizeof serv_addr.sin_zero); /* padding zeros*/

    //3. connect socket with corresponding address
    connect(sockfd, (struct sockaddr*)&serv_addr, sizeof(serv_addr));

    return sockfd;
}
```

```
struct sockaddr_in { // used for IPv4 only
    short sin_family; // addr family, AF_INET
    unsigned short sin_port; // port number
    struct in_addr sin_addr; // internet address
    unsigned char sin_zero[8]; // padding zeros
};
```

How to write a server: body (I)

```
int server_connect(unsigned int port_num)
/* port_num: which port to listen, return socket for subsequent communication*/
{
    int sockfd, new_fd;      /* listen on sock_fd, new connection on new_fd */
    struct sockaddr_in my_addr; /* my address */
    struct sockaddr_in their_addr; /* connector addr */
    int sin_size;

    /* create a socket */
    sockfd = socket(AF_INET, SOCK_STREAM, 0);
```

How to write a server: body (II)

```
struct sockaddr_in { // used for IPv4 only
    short    sin_family; // addr family, AF_INET
    unsigned short sin_port; // port number
    struct in_addr sin_addr; // internet address
    unsigned char sin_zero[8]; // padding zeros
};
```

```
// ...
/* set the address info */
my_addr.sin_family = AF_INET;
my_addr.sin_port = htons(port_num); /* short, network byte order */
my_addr.sin_addr.s_addr = INADDR_ANY
/* INADDR_ANY allows clients to connect to any one of the host's IP address.
```

```
memset(my_addr.sin_zero, '\0', sizeof(my_addr.sin_zero)); //padding zeros
```

```
/* bind the socket to the IP address and port number */
bind(sockfd, (struct sockaddr *) &my_addr, sizeof(struct sockaddr));
```

*The htons() function
converts from host byte order to network byte order*


*The inet_addr() function
converts the Internet host address from IPv4 numbers-and-dots notation
to network byte order.*

Functions

- **int listen(int sockfd, int backlog);**
 - Put socket into passive state (wait for connections rather than initiating a connection)
 - returns 0 on success, -1 and sets errno on failure
 - **sockfd**: socket file descriptor returned by socket()
 - **backlog**: the maximum number of connections this program can serve simultaneously

How to write a server: body (III)

```
// ...  
listen(sockfd, 5); /* maximum 5 pending connections */  
  
sin_size = sizeof(struct sockaddr_in);  
/* wait for client's connection, their_addr stores client's address */  
new_fd = accept(sockfd, (struct sockaddr*)&their_addr, &sin_size);  
  
return new_fd; /* new_fd is returned not sock_fd */  
}
```



store ip address of client

Functions

- **int accept(int sockfd, struct sockaddr* client_addr, int* addrlen);**
 - Accept a new connection
 - Return client's socket file descriptor or -1. Also sets errno on failure
 - **sockfd**: socket file descriptor for server, returned by socket()
 - **client_addr**: IP address and port number of a client (returned from call)
 - **addrlen**: length of address structure = pointer to **int** set to **sizeof(struct sockaddr_in)**
 - **NOTE: client_addr and addrlen are result arguments**
 - i.e. The program passes empty client_addr and addrlen into the function, and the kernel will fill in these arguments with client's information

Functions

- **int close(int sockfd);**
 - close a socket
 - return 0 on success, or -1 on failure
 - After close, sockfd is no longer valid

After the connection

For brevity, error handling is omit

//server side

```
int main()
{
    char buffer[100];
    int sock_fd = server_connect(80);
    write(sock_fd, "hahaha", 6);
    read(sock_fd, buffer, 100);
    //buffer will be "hihihi"
}
```

For brevity, error handling is omit

//client side

```
int main()
{
    char buffer[100];
    int sock_fd = client_connect("microsoft.com", 80);
    read(sock_fd, buffer, 100);
    //buffer will be "hahaha"
    write(sock_fd, "hihihi", 6);
}
```


Everything is a file descriptor: files, pipe, stdin, stdout, stderr, socket

Benefits: Use the same set of APIs: read, write, poll for everything.

Drawback: Implementation is very complex, loss of efficiency.

Client side code skeleton

```
int main(char * argc[], int argv) {  
    process_cmd_args(argc, argv);  
    socket_fd = client_connect(hostname, port);  
    //hostname, port number obtained from cmd args  
    set the terminal to the non-canonical, no echo mode  
    while (1) {  
        poll the socket_fd and stdin  
        if (socket_fd ready to read)  
            read from socket_fd, process special characters, send to stdout  
        if (stdin ready to read)  
            read from stdin, process special characters, send to stdout and socket_fd  
        if (error)  
            Proceed to exit process: read every last byte from socket_fd, write to stdout,  
            restore terminal and exit.  
    }  
}
```



Server side code skeleton

```
int main(char * argc[], int argv) {
```

```
    process_cmd_args(argc, argv);
```

```
    socket_fd = server_connect(port);
```

```
    //port number obtained from cmd args
```

```
    Register SIGPIPE handler, Create PIPE, fork the new process, perform stdin/stdout redirection,
```

```
    while (1) {
```

```
        poll the socket_fd and from_shell[0]
```

```
        if (socket_fd ready to read)
```

```
            read from socket_fd, process special characters, send to to_shell[1]
```

```
        if (from_shell[0] ready to read)
```

```
            read from from_shell[0], process special characters, send to socket_fd
```

```
        if (error)
```

```
            Proceed to exit process: read every last byte from from_shell[0], write  
            to socket_fd, get the exit status of the process and report to stderr.
```

```
    }
```

```
}
```

socket_fd
Socket

forward

server

Write

to_shell[1]

Read

from_shell[0]

shell

Read

Write

(stdout/stderr)



Compress

- Purpose
 - Reduce the amount of space data takes up (less bandwidth)
 - Require processing at sending/receiving end
 - Tradeoff between bandwidth and processing cost
- Use standard compression library
 - Depend on your system, install [zilib](#) package (library, documentation)
- Start simple
 - A simple program that reads from and writes to a pipe.
- Add **--compress** option (client & server)
 - If included, will enable compression (of all traffic in [both directions](#)).
 - compress traffic *before sending* it over the network
 - decompress traffic *after receiving* it


Compress

- Use **--log** option to verify that compression/decompression works
 - You should record
 - outgoing data post-compression
 - incoming data pre-decompression
 - You should expect to see:
 - no clear text going either to or from the shell
 - a reduction in the number of bytes sent in both directions

Compress the network connections Overview

Achieved with the zlib library.

```
int main(void)
{
    z_stream stream;
    init_compress_stream (&stream); //init a compress stream.
    while (!done) {
        do_compress (&stream, orig_buf, orig_len, out_buf, out_len);
        //Use the zlib to compress the data from orig_buf and put it in out_buf.
        output the out_buf to the network.
    }
    end_compress_stream(&stream); //Close the compress stream
}
```



Note: purple text are user defined function names to encapsulate certain logic of the code, it's not standard functions from the zlib library

Basic functions

- From `zlib.h` (`#include "zlib.h"`)
 - basic compression functions
 - `deflateInit()`, `deflate()`, and `deflateEnd()`
 - basic decompression functions
 - `inflateInit()`, `inflate()`, and `inflateEnd()`
 - stream data structure
 - The `strm` structure is used to pass information to and from the *zlib* routines

Compress

- Overall Logic

```
z_stream strm; // strm structure is used to pass information to and from the zlib routines, and to  
maintain the deflate() state
```

```
unsigned char inbuf[CHUNK]; // input buffer for deflate().
```

```
unsigned char outbuf[CHUNK]; // output buffer for deflate().
```

```
//initialize compression stream
```

```
//while loop to compress data from inputbut to outputbuf
```

```
//end compression stream
```


initialize the zlib state for compression using **deflateInit()**

- The **zalloc**, **zfree**, and **opaque** fields in the strm structure must be initialized before calling deflateInit()

```
/* allocate deflate state */
```

```
strm.zalloc = Z_NULL; // set to Z_NULL, to request zlib use the default memory allocation routines
```

```
strm.zfree = Z_NULL;
```

```
strm.opaque = Z_NULL;
```

```
ret = deflateInit(&strm, Z_DEFAULT_COMPRESSION); /* pointer to the structure to be initialized, and the  
compression level */
```

```
if (ret != Z_OK) // make sure that it was able to allocate memory, and the provided arguments were valid.
```

```
{
```

```
    // code here: error handling
```

```
}
```

```
ret = deflateInit(&strm, level);
```

- **strm**: pointer to the structure to be initialized and
- **level**: the compression level (integer in the range of -1 to 9)
 - Lower compression levels result in faster execution, but less compression. Higher levels result in greater compression, but slower execution.
 - The zlib constant Z_DEFAULT_COMPRESSION, equal to -1, provides a good compromise between compression and speed and is equivalent to level 6.
- Return value **ret**:
 - Check against the *zlib* constant Z_OK

Deflate()

```
strm.avail_in = /* number of bytes available at next_in */
strm.next_in = inbuf; /* next input byte*/
strm.avail_out = sizeof outbuf; /* remaining free space at next_out */
strm.next_out = outbuf; /* next output byte */

while (strm.avail_in > 0) {
    deflate(&strm, Z_SYNC_FLUSH); //use Z_SYNC_FLUSH for independent msgs
    /* deflate(). It takes as many of the avail_in bytes at next_in as it can process, and writes
    as many as avail_out bytes to next_out.
    deflate will update next_in, avail_in, next_out, avail_out accordingly.*/
}
```

Deflate()

```
strm.avail_in = /* number of bytes available at next_in */ minus 100bytes
strm.next_in = inbuf; /* next input byte*/ minus 100bytes of data
strm.avail_out = sizeof outbuf; /* remaining free space at next_out */ decrease
strm.next_out = outbuf; /* next output byte */ increase

while (strm.avail_in > 0) {
    deflate(&strm, Z_SYNC_FLUSH); //use Z_SYNC_FLUSH for independent msgs
    Compress 100bytes of data
    deflate will update next_in, avail_in, next_out, avail_out accordingly.*
    compressed_bytes = sizeof outbuf - strm.avail_out; //calculate size of compressed data
}
```

Close the module

```
deflateEnd(&stream); //frees up data structure
```

Ref:

<https://www.zlib.net/manual.html>

https://www.zlib.net/zlib_how.html