

CS 111

Operating Systems Principles

Section 1E Week 3

Tengyu Liu


Changes

- Slides are now uploaded to CCLE every Wednesday
 - I will update the slides before each discussion section
- I will explain the overall picture of each project
- I will include more code examples


Overview

- Project 1B: Inter-Process Communication over Internet
- System calls:
 - Internet communication
 - `socket(2)`, `connect(2)`, `bind(2)`, `listen(2)`, `accept(2)`, `shutdown(2)`
 - Compressed Communication
 - `deflateInit()`, `inflateInit()`, `deflate()`, `inflate()`, `deflateEnd()`, `inflateEnd()`

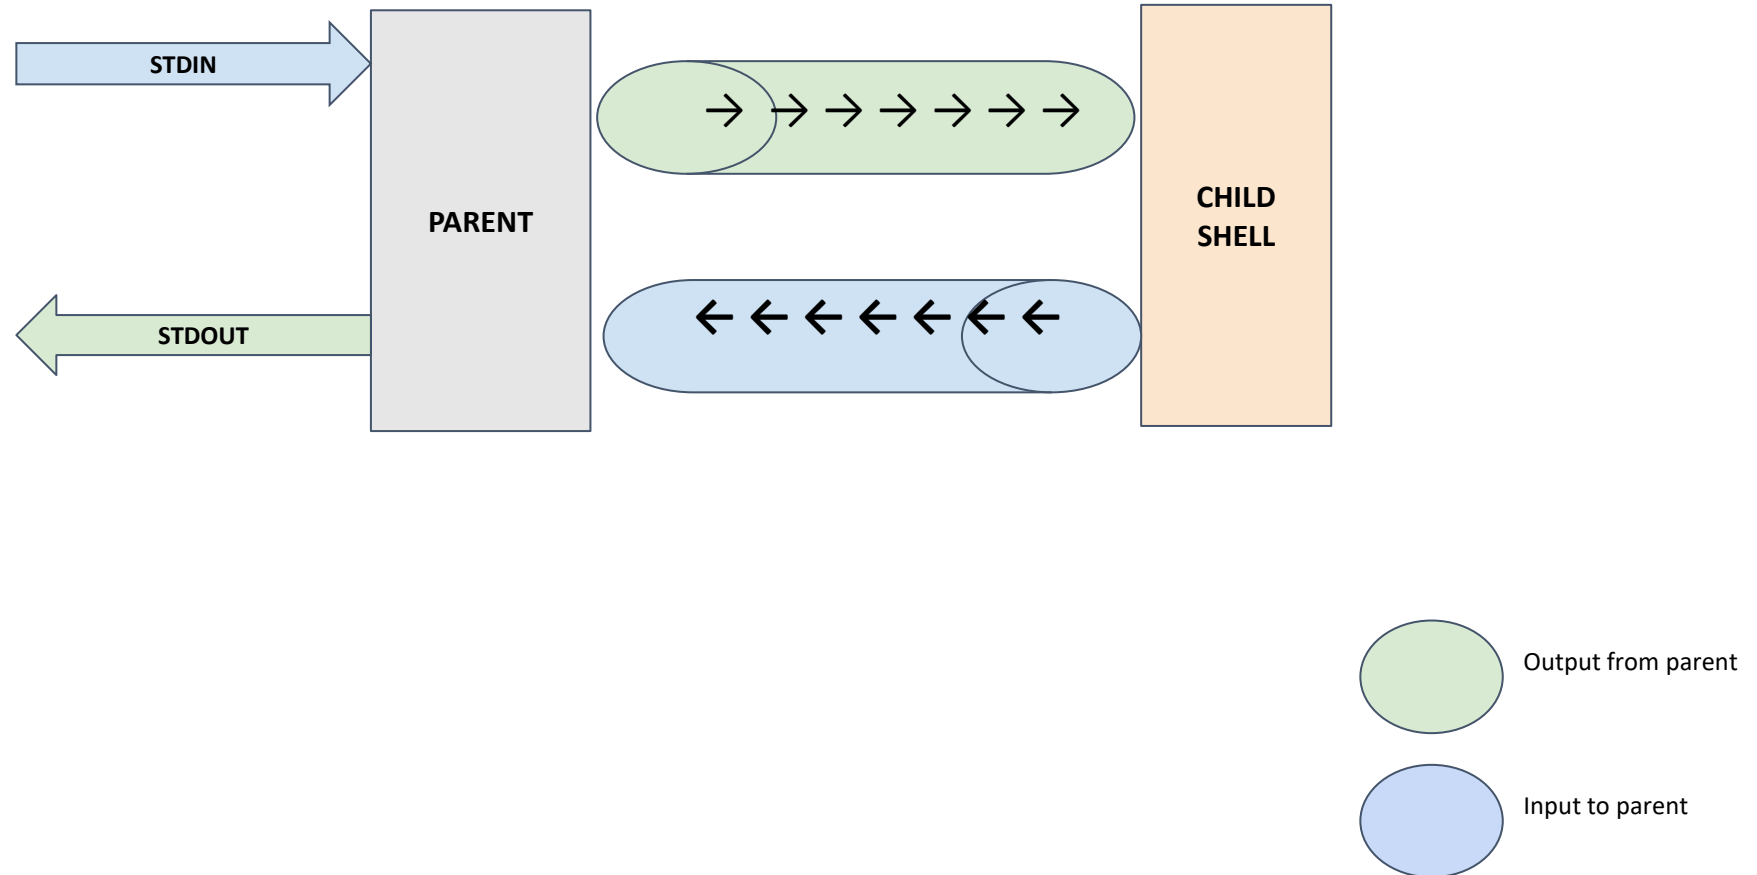
Project 1B




Project 1B



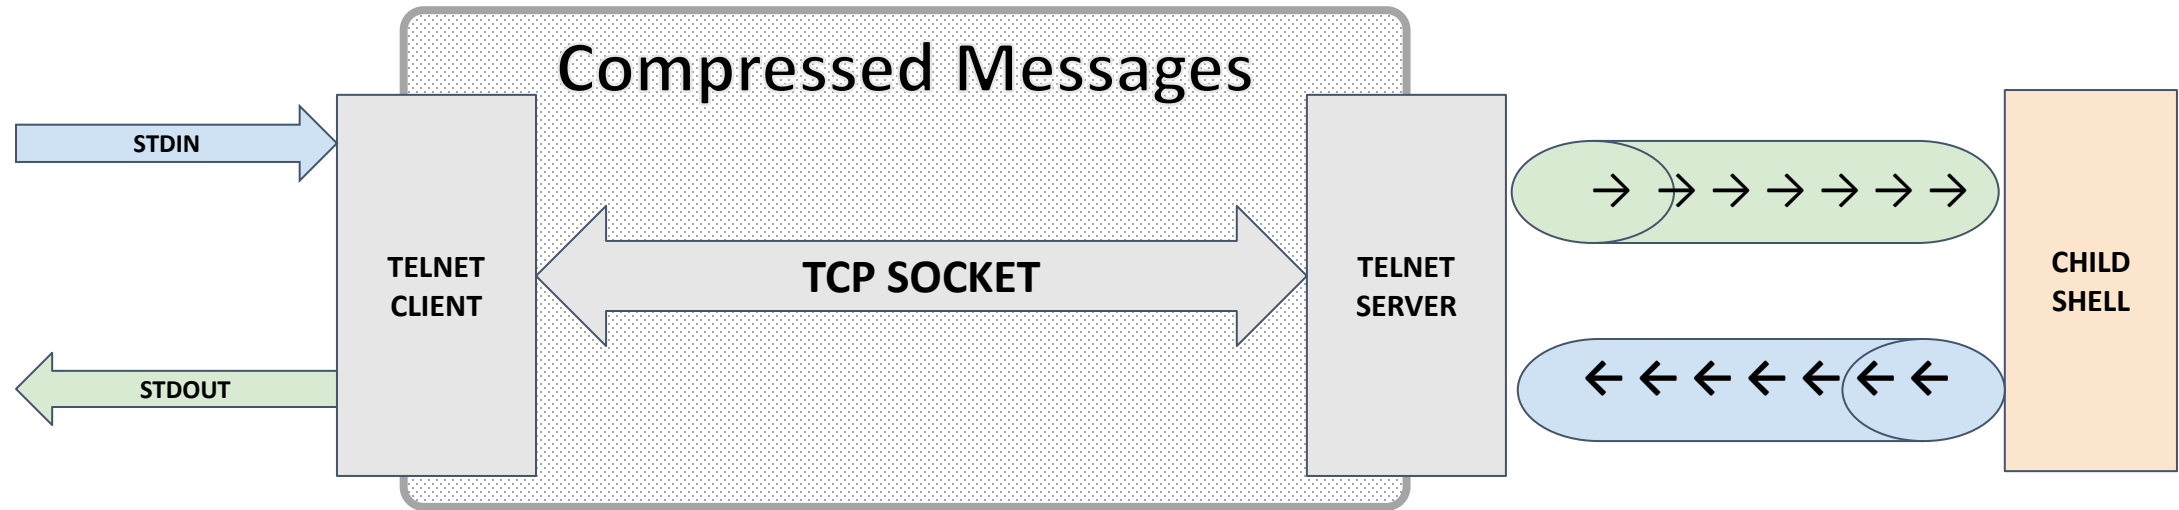
Project 1A



Project 1B



Project 1B --compress



TCP Socket

TCP Socket

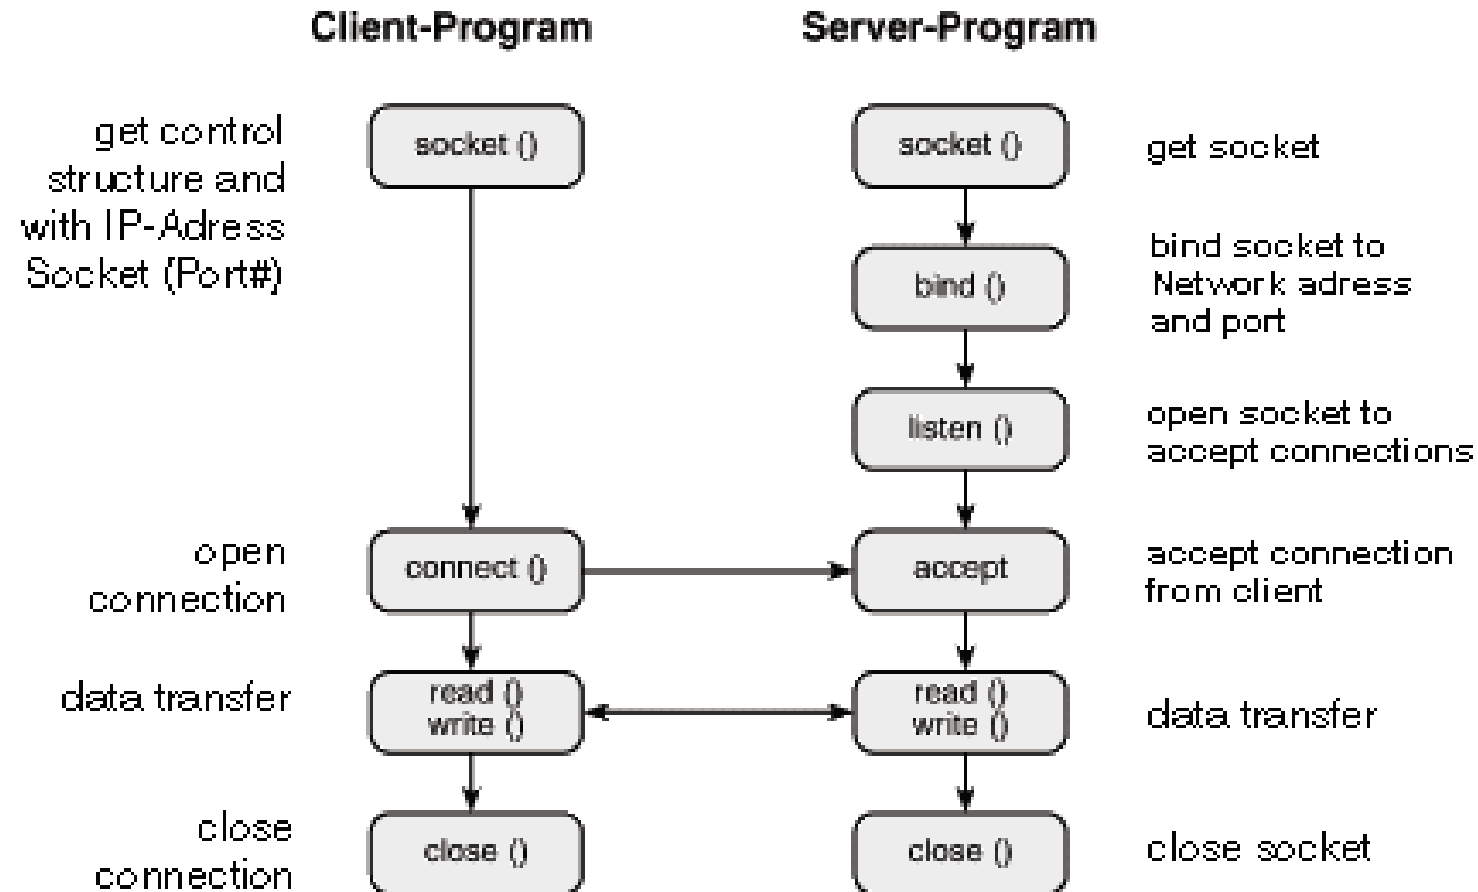
Pipe

- Both ends must live on the same machine
- Returns 2 file descriptors, one for reading, one for writing

Socket

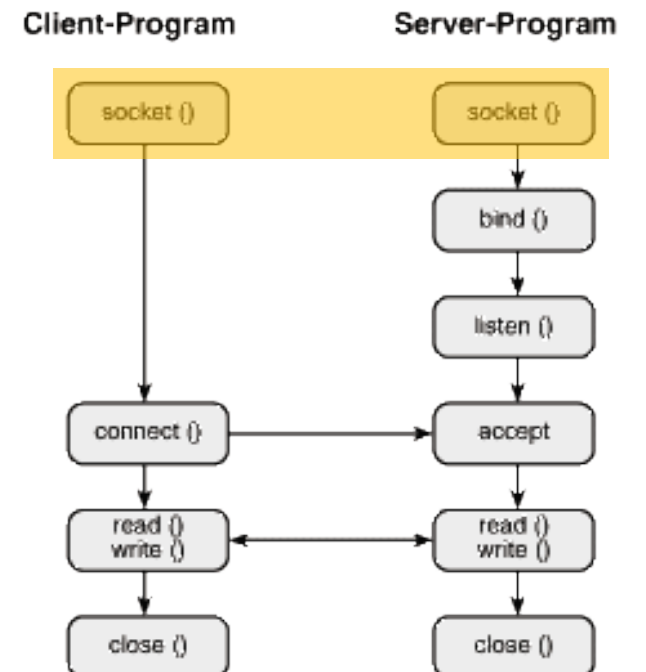
- Makes no assumption on process location
- Returns only 1 file descriptor, for both reading and writing
- Socket needs more information
 - Protocol
 - Target address

TCP Socket: Pipeline Overview




socket(2)

- Create an endpoint for communication
- `int socket(int domain, int type, int protocol)`




socket(2)

- Create an endpoint for communication
- `int socket(int domain, int type, int protocol)`
 - Communication domain
 - IPv4: `AF_INET`
 - IPv6: `AF_INET6`
 - Local: `AF_LOCAL`



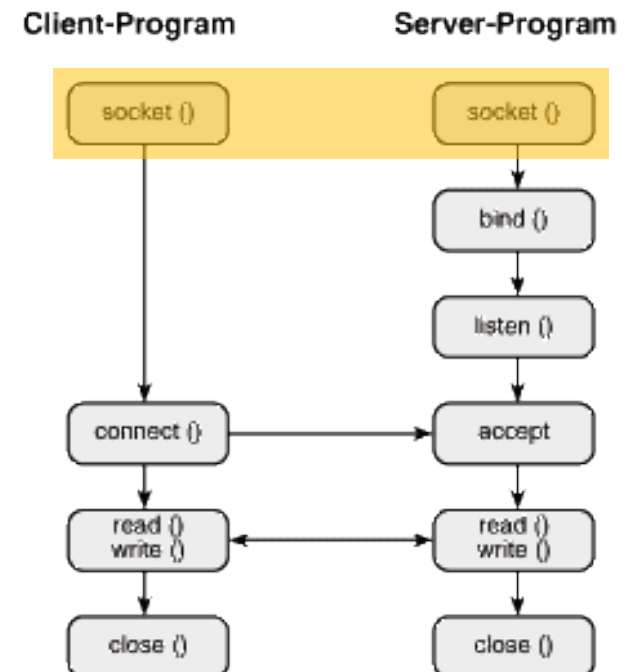
socket(2)

- Create an endpoint for communication
- `int socket(int domain, int type, int protocol)`
 - Communication type
 - TCP: `SOCK_STREAM`
 - UDP: `SOCK_DGRAM`




socket(2)

- Create an endpoint for communication
- `int socket(int domain, int type, int protocol)`
 - IP protocol value
 - Usually there is only one protocol available for each domain
 - Use 0 to use the default



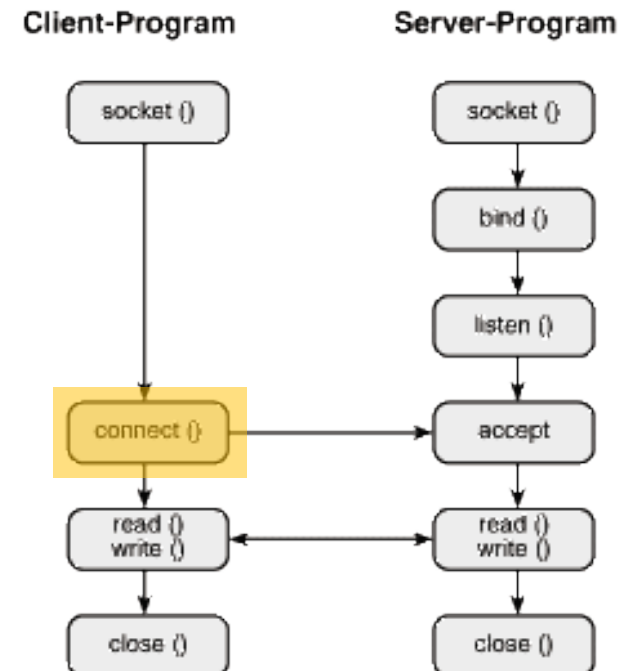
socket(2)

- Create an endpoint for communication
- **int socket(int domain, int type, int protocol)**
 - Return value: socket descriptor
 - On error: returns -1



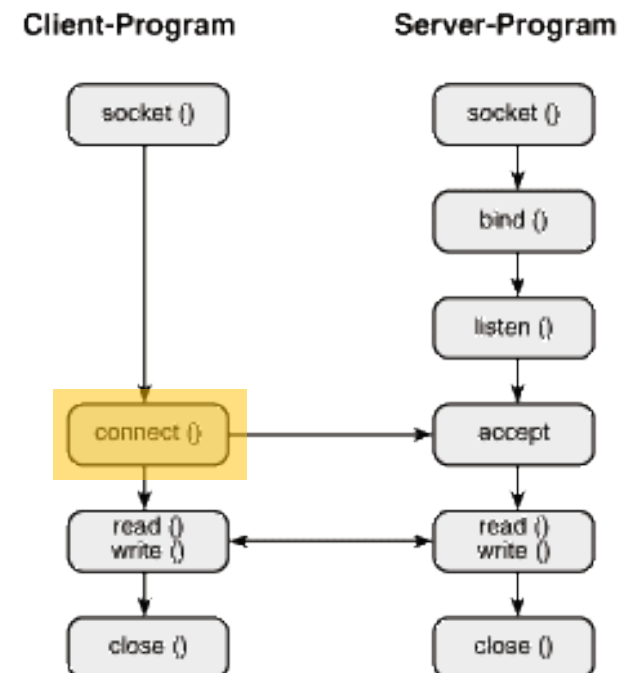
connect(2)

- Connect socket to a remote host
- `int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);`



connect(2)

- Connect socket to a remote host
- `int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);`
 - Socket for connect (returned by `socket(2)`)

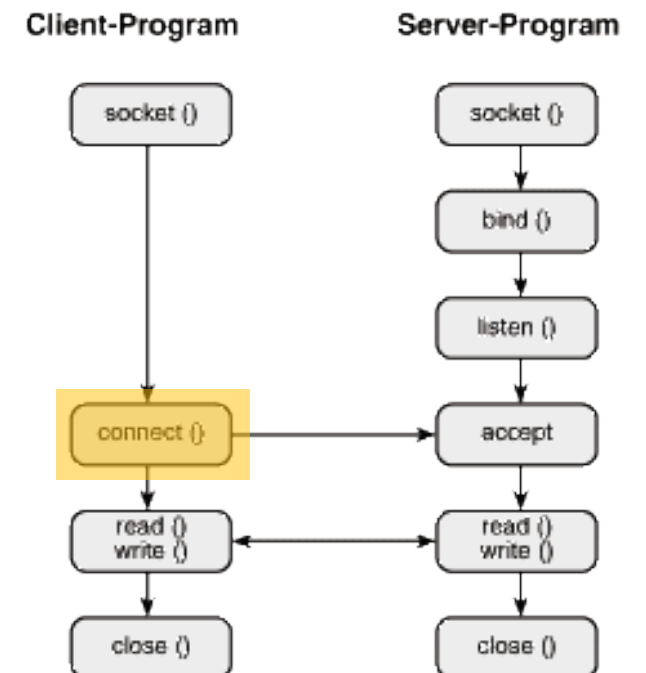


connect(2)

- Connect socket to a remote host
- `int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);`
 - Structure containing server's IP address and port number

```
struct sockaddr_in {
    sa_family_t    sin_family; /* address family: AF_INET */
    in_port_t      sin_port;   /* port in network byte order */
    struct in_addr sin_addr;    /* internet address */
};

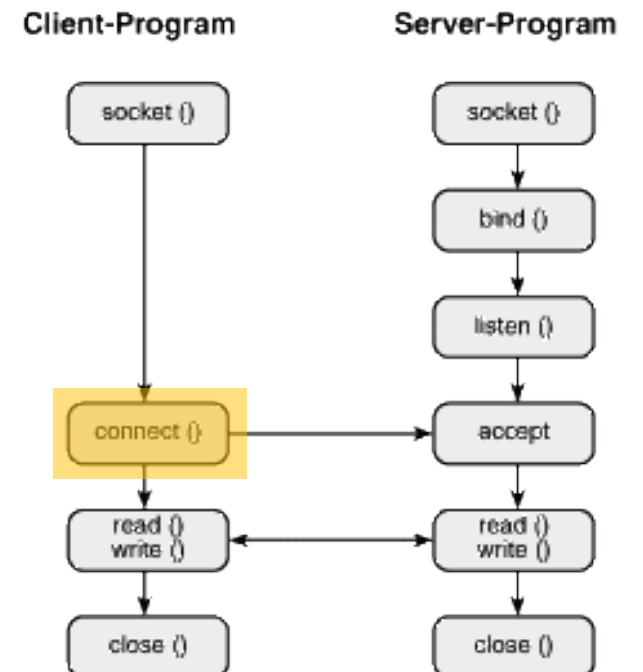
/* Internet address. */
struct in_addr {
    uint32_t        s_addr;    /* address in network byte order */
};
```



connect(2)

- Connect socket to a remote host
- `int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);`
 - Structure containing server's IP address and port number

```
struct sockaddr_in {  
    AF_INET → sa_family_t    sin_family; /* address family: AF_INET */  
              in_port_t      sin_port;   /* port in network byte order */  
              struct in_addr sin_addr;   /* internet address */  
};  
  
/* Internet address. */  
struct in_addr {  
    uint32_t      s_addr;   /* address in network byte order */  
};
```

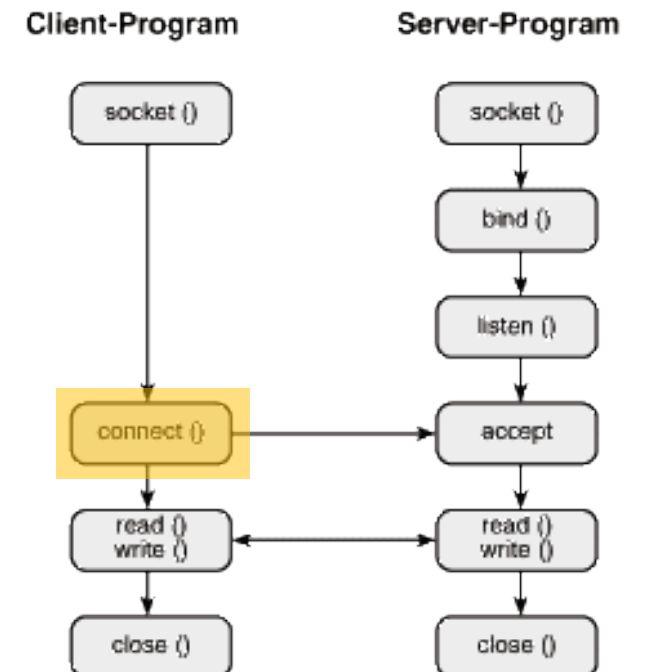


connect(2)

- Connect socket to a remote host
- `int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);`
 - Structure containing server's IP address and port number

```
struct sockaddr_in {
    sa_family_t    sin_family; /* address family: AF_INET */
    Port number → in_port_t    sin_port; /* port in network byte order */
    struct in_addr sin_addr; /* internet address */
};

/* Internet address. */
struct in_addr {
    uint32_t        s_addr; /* address in network byte order */
};
```

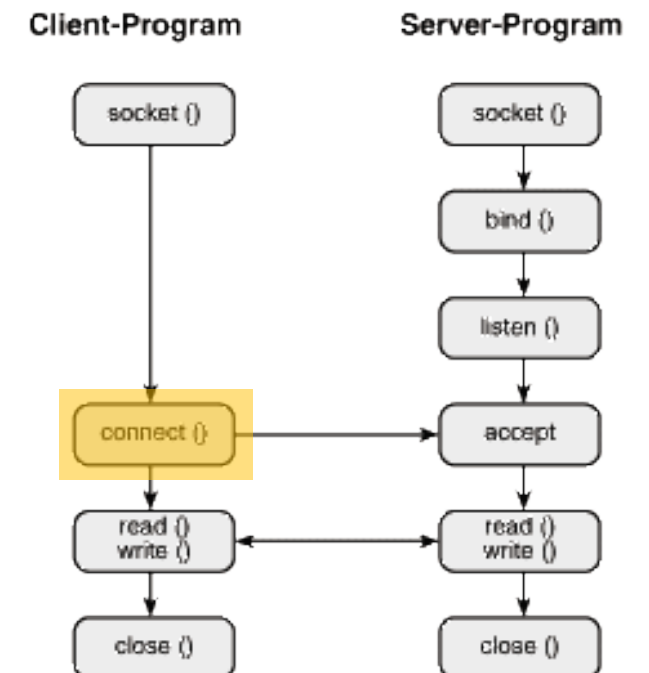


connect(2)

- Connect socket to a remote host
- `int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);`
 - Structure containing server's IP address and port number


```
struct sockaddr_in {
    sa_family_t    sin_family; /* address family: AF_INET */
    in_port_t      sin_port;   /* port in network byte order */
    struct in_addr sin_addr;    /* internet address */
};

/* Internet address. */
struct in_addr {
    Obtained by
    gethostbyname(3) → uint32_t    s_addr; /* address in network byte order */
};
```




connect(2)

- Connect socket to a remote host
- `int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);`
 - `sizeof(*addr)`



connect(2)

- Connect socket to a remote host
- **int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);**
 - Returns the socket descriptor
 - On error, return -1



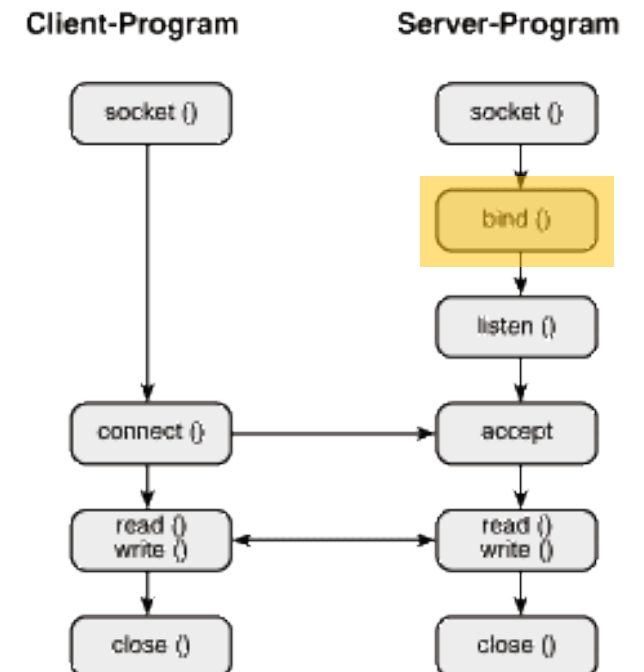
connect(2)

- `int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);`
- Example (error checks are omitted)

```
// get target ip address
char *host = "localhost";
struct hostent *server = gethostbyname(host);
// construct sockaddr_in struct
struct sockaddr_in serv_addr;
bzero( (char *) &serv_addr, sizeof(serv_addr));
bcopy((char *) server->h_addr, (char *) &serv_addr.sin_addr.s_addr, server->h_length);
serv_addr.sin_family = AF_INET;      // specify IPv4
serv_addr.sin_port = htons(port);    // specify port number
connect(sockfd, (struct sockaddr *) &serv_addr, sizeof(serv_addr))
```


bind(2)

- Bind a name to a socket
- `int bind(int sockfd, const struct sockaddr *addr, socklen_t addrlen);`



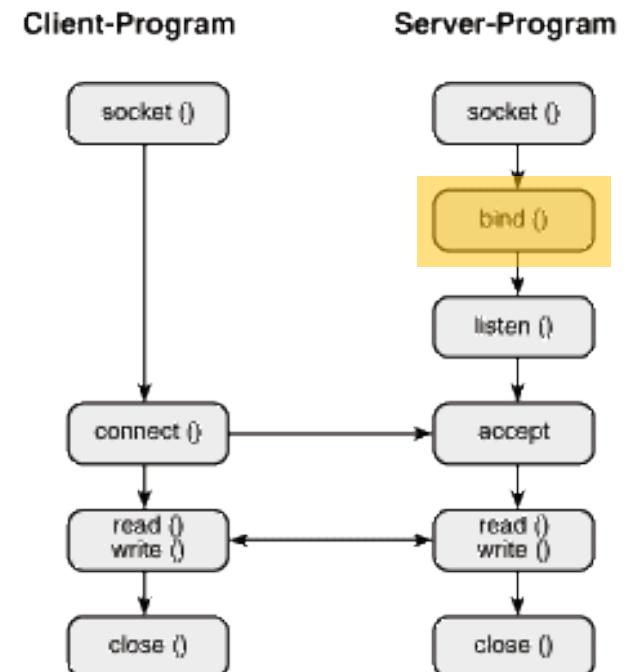
bind(2)

- Bind a name to a socket
- `int bind(int sockfd, const struct sockaddr *addr, socklen_t addrlen);`
 - Socket descriptor




bind(2)

- Bind a name to a socket
- `int bind(int sockfd, const struct sockaddr *addr, socklen_t addrlen);`
 - IP Address and Port




bind(2)

- Bind a name to a socket
- `int bind(int sockfd, const struct sockaddr *addr, socklen_t addrlen);`
 - `sizeof(*addr)`



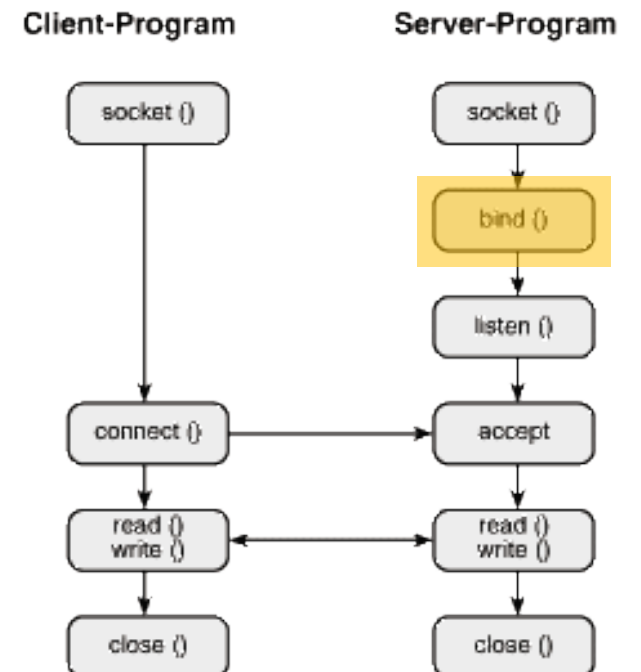
bind(2)

- Bind a name to a socket
- **int bind(int sockfd, const struct sockaddr *addr, socklen_t addrlen);**
 - Socket descriptor
 - On error, return -1



bind(2)

- Bind a name to a socket
- `int bind(int sockfd, const struct sockaddr *addr, socklen_t addrlen);`
 - This function effectively gives a port to the socket
 - IP address: address of the machine
 - Port number: address of the process
 - Without a port number, messages cannot reach your process

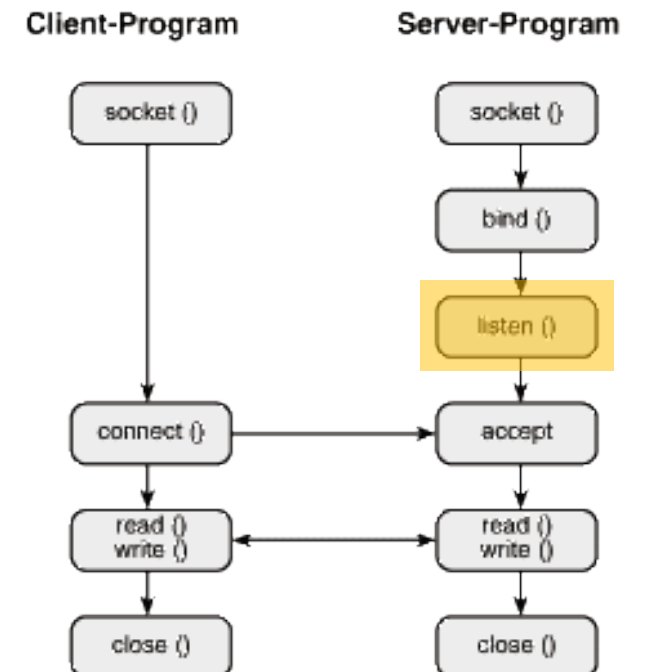


bind(2)

```
listenfd = socket(AF_INET, SOCK_STREAM, 0);  
bzero((char *) &serv_addr, sizeof(serv_addr));  
serv_addr.sin_family = AF_INET;  
serv_addr.sin_addr.s_addr = INADDR_ANY;  
serv_addr.sin_port = htons(port);  
bind(listenfd, (struct sockaddr *) &serv_addr, sizeof(serv_addr))
```

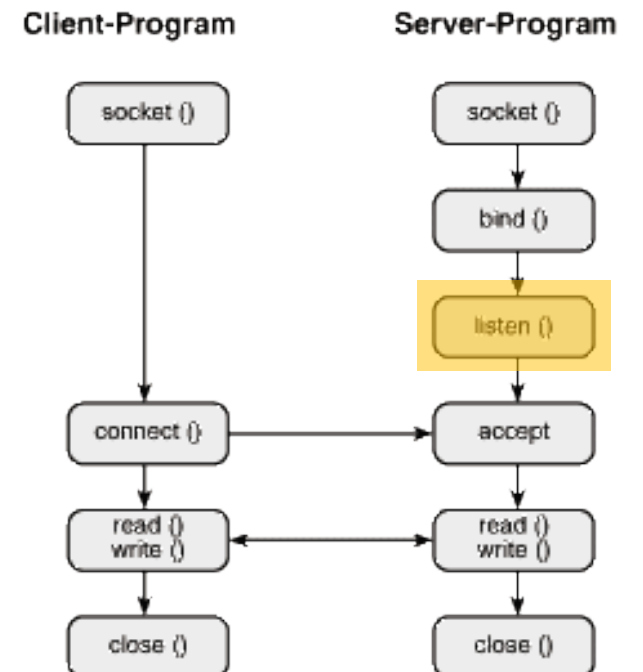

listen(2)

- Start listening for incoming connections
- `int listen(int sockfd, int backlog)`



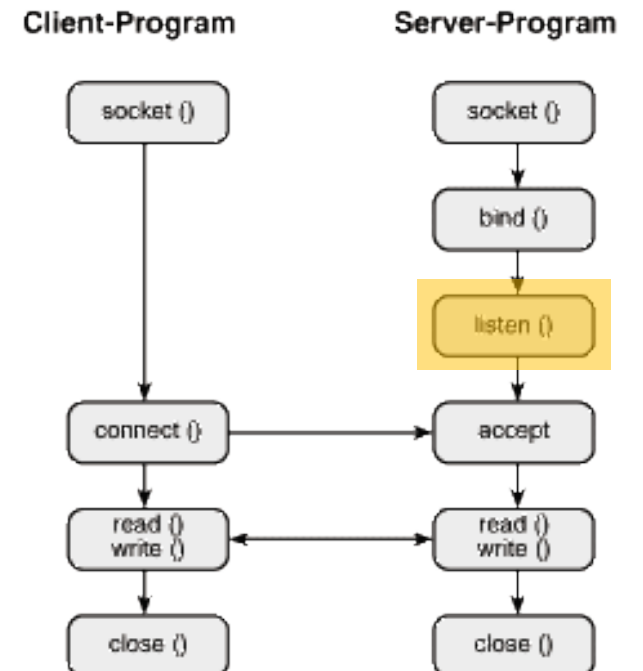
listen(2)

- Start listening for incoming connections
- `int listen(int sockfd, int backlog)`
 - Socket file descriptor



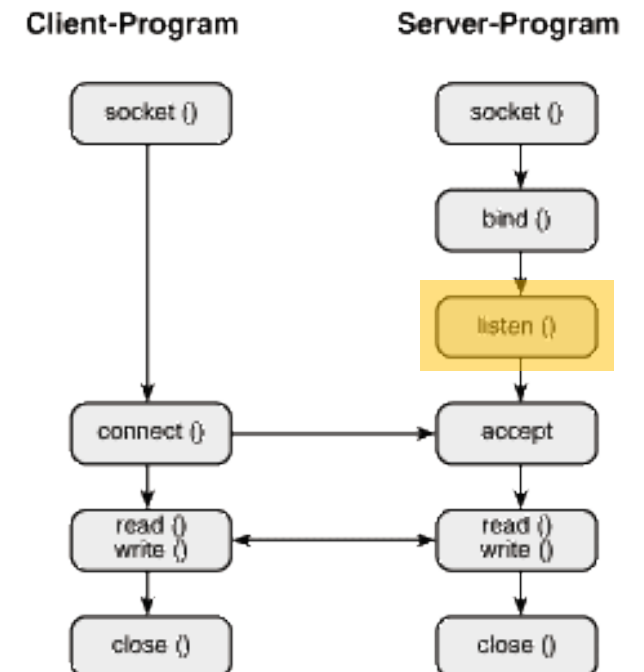
listen(2)

- Start listening for incoming connections
- `int listen(int sockfd, int backlog)`
 - Number of connections allowed in the incoming queue
 - On most systems, max allowed is 5
 - Use 5




listen(2)

- Start listening for incoming connections
- **int listen(int sockfd, int backlog)**
 - Success: 0
 - Error: 1




accept(2)

- Accept a connection on a socket (blocking until a connection is accepted)
- `int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen)`




accept(2)

- Accept a connection on a socket (blocking until a connection is accepted)
- `int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen)`
 - Socket descriptor that we listened on



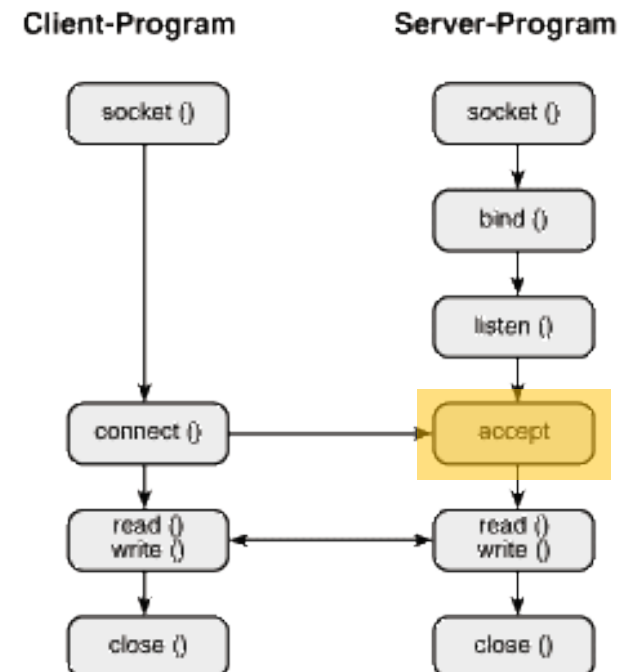
accept(2)

- Accept a connection on a socket (blocking until a connection is accepted)
- `int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen)`
 - Output parameter
 - Contains information about the incoming connection



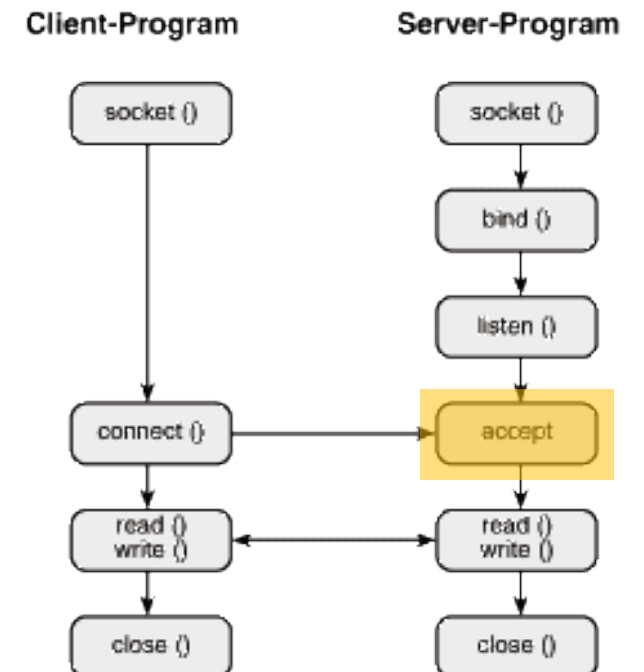
accept(2)

- Accept a connection on a socket (blocking until a connection is accepted)
- `int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen)`
 - Length of the incoming `sockaddr` struct



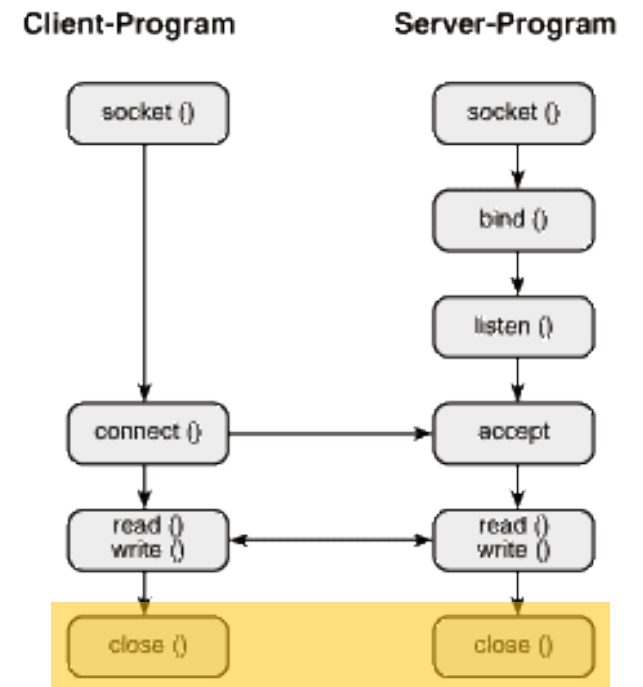
accept(2)

- Accept a connection on a socket (blocking until a connection is accepted)
- **int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen)**
 - Success
 - New socket file descriptor to use for this connection
 - The old descriptor is only for accepting incoming connections
 - A new descriptor is allocated for every connection established
 - Error
 - -1




shutdown(2)

- Shut down socket
- `int shutdown(int socket, int how)`



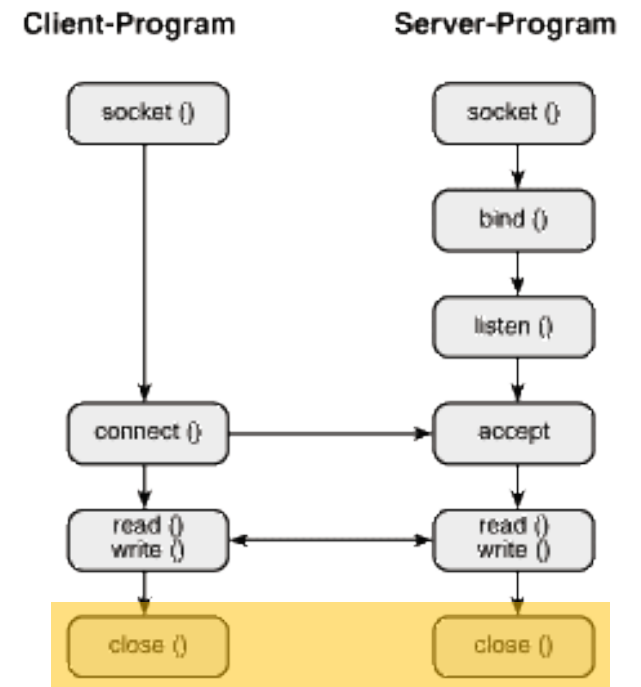
shutdown(2)

- Shut down socket
- `int shutdown(int socket, int how)`
 - Socket descriptor




shutdown(2)

- Shut down socket
- `int shutdown(int socket, int how)`
 - SHUT_RD: shut down receive ops
 - SHUT_WR: shut down sending ops
 - SHUT_RDWR: shut down obth ops



shutdown(2)

- Shut down socket
- **int shutdown(int socket, int how)**
 - Success: 0
 - Error: -1



Compressed Communication

Advantage?

Disadvantage?

Compressed Communication

Advantage

- Reduce communication length
- Therefore uses less bandwidth

Disadvantage

- Requires processing at both ends of the communication channel

Compressed Communication

Advantage

- Reduce communication length
- Therefore uses less bandwidth

Disadvantage

- Requires processing at both ends of the communication channel

Is compressed message always smaller than the uncompressed one?

Can compression provide privacy?

Compressed communication

- Pipeline
 - To compress
 - `deflateInit()`, `deflate()`, `deflateEnd()`
 - To decompress
 - `inflateInit()`, `inflate()`, `inflateEnd()`
- You will need zlib
 - `include <zlib.h>`
 - Compile with `-lz` flag

deflateInit()

- Initialize internal stream state for compression
- `ZEXTERN int ZEXPORT deflateInit OF((z_streamp strm, int level))`

deflateInit()

- Initialize internal stream state for compression
- **ZEXTERN int ZEXPORT deflateInit OF((z_streamp strm, int level))**
 - zlib macros for cross-platform compatibility
 - You can think of it as
 - `extern int deflateInit(z_streamp strm, int level)`

deflateInit()

- Initialize internal stream state for compression
- ZEXTERN int ZEXPORT deflateInit OF((z_streamp strm, int level))
 - Pointer to z_stream

initialize to Z_NULL

```
typedef struct z_stream_s {
    z_const Bytef *next_in;      /* next input byte */
    uInt      avail_in;  /* number of bytes available at next_in */
    uLong     total_in;  /* total number of input bytes read so far */

    Bytef      *next_out; /* next output byte will go here */
    uInt      avail_out; /* remaining free space at next_out */
    uLong     total_out; /* total number of bytes output so far */

    z_const char *msg; /* last error message, NULL if no error */
    struct internal_state FAR *state; /* not visible by applications */

    alloc_func zalloc; /* used to allocate the internal state */
    free_func  zfree;  /* used to free the internal state */
    voidpf     opaque; /* private data object passed to zalloc and zfree */

    int      data_type; /* best guess about the data type: binary or text
                        for deflate, or the decoding state for inflate */
    uLong     Adler;     /* Adler-32 or CRC-32 value of the uncompressed data */
    uLong     reserved; /* reserved for future use */
} z_stream;
```

deflateInit()

- Initialize internal stream state for compression
- `ZEXTERN int ZEXPORT deflateInit OF((z_streamp strm, int level))`
 - Between 0 (no compression) and 9 (most compression)
 - Can set to `Z_DEFAULT_COMPRESSION`

deflateInit()

- Initialize internal stream state for compression
- **ZEXTERN int ZEXPORT deflateInit** OF((z_streamp strm, int level))

```
#define Z_OK          0
#define Z_STREAM_END  1
#define Z_NEED_DICT   2
#define Z_ERRNO       (-1)
#define Z_STREAM_ERROR (-2)
#define Z_DATA_ERROR  (-3)
#define Z_MEM_ERROR   (-4)
#define Z_BUF_ERROR   (-5)
#define Z_VERSION_ERROR (-6)
```

inflateInit()

- Initialize streams for decompression
- `ZEXTERN int ZEXPORT inflateInit OF((z_streamp strm))`

inflateInit()

- Initialize streams for decompression
- `ZEXTERN int ZEXPORT inflateInit OF((z_streamp strm))`
 - Pointer to `z_stream`
 - Same as `deflateInit()`

inflateInit()

- Initialize streams for decompression
- `ZEXTERN int ZEXPORT inflateInit OF((z_streamp strm))`
 - Same as `deflateInit()`

deflate()

- Compress data until input buffer is empty or output buffer is full
- ZEXTERN int ZEXPORT deflate OF((z_streamp strm, int flush))

deflate()

- Compress data until input buffer is empty or output buffer is full
- ZEXTERN int ZEXPORT deflate OF((z_streamp strm, int flush))
 - Same as before
 - strm.next_in: next input byte
 - strm.avail_in: number of bytes available in next_in
 - strm.total_in: total number of bytes read so far
 - Similarly,
 - strm.next_out, strm.avail_out, strm.total_out

deflate()

- Compress data until input buffer is empty or output buffer is full
- ZEXTERN int ZEXPORT deflate OF((z_streamp strm, **int flush**))
 - How do you want to force flush
 - Forcing flush frequently degrades compression ratio
 - Z_NO_FLUSH gives best compression ratio

```
#define Z_NO_FLUSH      0
#define Z_PARTIAL_FLUSH 1
#define Z_SYNC_FLUSH    2
#define Z_FULL_FLUSH    3
#define Z_FINISH        4
#define Z_BLOCK         5
#define Z_TREES         6
```

deflate()

- Compress data until input buffer is empty or output buffer is full
- ZEXTERN **int** ZEXPORT **deflate** OF((z_streamp strm, int flush))
 - Z_OK on success

inflate()

- Decompress data until input buffer is empty or output buffer is full
- ZEXTERN int ZEXPORT inflate OF((z_streamp strm, int flush))
 - Similar to deflate()

deflateEnd() / inflateEnd()

- ZEXTERN int ZEXPORT deflateEnd OF((z_streamp strm))
- ZEXTERN int ZEXPORT inflateEnd OF((z_streamp strm))

An example code for compressing

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
#include <zlib.h>

int main( int argc, char **argv ) {
    // initialize compressor
    z_stream compressor;
    // message to be compressed
    const char* input_string = "abcdabcdabdefgefgefgef";
    // initialize compressor
    compressor.zalloc = Z_NULL;
    compressor.zfree = Z_NULL;
    compressor.opaque = Z_NULL;
    compressor.avail_in = 0;
    compressor.next_in = Z_NULL;
    int ret = deflateInit(&compressor, Z_DEFAULT_COMPRESSION);
    if (ret != Z_OK) {
        exit(1);
    }

    // prepare for compression
    char output_buf[1024];
    compressor.avail_in = sizeof input_string;
    compressor.next_in = (unsigned char *) input_string;
    compressor.avail_out = sizeof output_buf;
    compressor.next_out = (unsigned char *) output_buf;
    // compress message
    do {
        (void) deflate(&compressor, Z_SYNC_FLUSH);
    } while( compressor.avail_in > 0 );
    // print the compressed message
    write(1, output_buf, sizeof output_buf - compressor.avail_out);
    write(1, "\n", 1);
    // shutdown compressor
    deflateEnd(&compressor);
}
```

```
> ./compress
xJLJNl@b
```


Workflow

- How your code should work

Workflow: client

- Initialize zlib streams
- Create a socket
 - `socket(2)`
- Identify server
 - `gethostbyname(3)`
- `connect(2)` to server
- Wait for input
 - `poll(2)` on keyboard and socket
 - Compress and decompress if necessary
- `shutdown(2)` socket
- Restore terminal modes

Workflow: server

- Initialize zlib streams
- Create a socket
 - `socket(2)`
- `bind(2)` socket to name
 - Fill server's `sockaddr_in` struct `INADDR_ANY`
- Establish connection with client
 - `listen(2)`, `accept(2)`
- Input/output forwarding
 - `poll(2)` on socket and shell, decompress if necessary

Recommendation

- Test individual modules before putting everything together
 1. Write a server and a client that simply sends “hello” to each other
 2. Write a program and compresses and decompresses messages
 - Test if you compress and decompress, you get the same message back
 3. Add --compress support to your server and client program
 4. Add project 1A to your program

Socket FAQ

- Socket-in-use error
 - `listen(2)` remains active for some time (usually seconds) after the server exits. Wait for a few seconds and try again. Or you can switch to another port.
- Connection refused error
 - Nothing is listening to the port you are trying to connect.
 - On linux servers, monitor listening ports with
 - `netstat -tln`