

CS 111

Operating Systems Principles

Section 1E Week 5

Tengyu Liu

# Project 4

- If you don't have a beaglebone kit, **BUY ONE NOW!**

# Previously in project 2A...

- Time per operation increases with the number of threads
  - Shouldn't it be decreasing?

# Previously in project 2A...

- Time per operation increases with the number of threads
  - Shouldn't it be decreasing?
- The problem
  - A lot of waiting during each operation
  - While waiting, other threads are working in parallel
  - More threads = more waiting

# Previously in project 2A...

- Time per operation increases with the number of threads
  - Shouldn't it be decreasing?
- The problem
  - A lot of waiting during each operation
  - While waiting, other threads are working in parallel
  - More threads = more waiting
- The bottleneck
  - Parallel threads cannot run in parallel if they acquire the same lock

# Previously in project 2A...

- Multithreading does slow down things...
  - All threads can run in parallel until they acquire a lock
  - At that point they have to wait in line
  - This is called **contention**

# Previously in project 2A...

- Multithreading does slow down things...
  - All threads can run in parallel until they acquire a lock
  - At that point they have to wait in line
  - This is called **contention**
- Contention is worse for list operations. Why?

# Previously in project 2A...

- Multithreading does slow down things...
  - All threads can run in parallel until they acquire a lock
  - At that point they have to wait in line
  - This is called **contention**
- Contention is worse for list operations. Why?
  - Inserting in a list requires more work than addition
  - The time spent in the critical section is higher for list operations
  - Therefore all other threads need to wait longer



# Previously in project 2A...

- Multithreading does slow down things...
  - All threads can run in parallel until they acquire a lock
  - At that point they have to wait in line
  - This is called **contention**
- Contention is worse for list operations. Why?
  - Inserting in a list requires more work than addition
  - The time spent in the critical section is higher for list operations
  - Therefore all other threads need to wait longer
- How can we decrease contention for lists?

# Previously in project 2A...

- Multithreading does slow down things...
  - All threads can run in parallel until they acquire a lock
  - At that point they have to wait in line
  - This is called **contention**
- Contention is worse for list operations. Why?
  - Inserting in a list requires more work than addition
  - The time spent in the critical section is higher for list operations
  - Therefore all other threads need to wait longer
- How can we decrease contention for lists?
  - Splitting lists into sub-lists and lock them independently

# Project 2B

## Part A

- Identify the contention problem
  - Starting from project 2A, measure the throughput (# of total ops/sec)
  - Pinpoint the most time-consuming part of the program

## Part B

- Resolve the lock contention problem

# Measuring throughput

- In a single-threaded program, time per op is a reasonable performance measure
- For multi-threaded implementations, we are also concerned about **how well we are taking advantage of parallelism**
- Measuring the aggregated throughput
  - # of total ops / total time taken
  - Plot with 1000 iterations and (1,2,4,8,12,16,24) threads for both mutex and spin-lock
  - You should see the throughput drop as #threads increase

# Pinpointing bottleneck

- You used 2 types of locks to protect critical sections
  - Spin-locks
  - Mutexes
- Both prevented race conditions
  - Do they add any synchronization cost?
  - Is one better than another?

# Synchronization cost

**Spin-Lock**

**Mutex**

# Synchronization cost

## **Spin-Lock**

- CPU cycles are used as thread tries to acquire the lock

## **Mutex**

- Context switches when the thread can't acquire the lock

# Synchronization cost

## **Spin-Lock**

- CPU cycles are used as thread tries to acquire the lock

## **Mutex**

- Context switches when the thread can't acquire the lock

How can we measure that overhead?



# Measuring the cost for spin-locks

- CPU Profiling tool
  - Takes a snapshot of the CPU every millisecond
    - Interval is a parameter
  - At each snapshot, it records what function is running
  - After the process has finished, we can look at the log
  - If 100 snapshots were taken in function X, we can estimate that function X ran for 100ms

# Measuring the overhead for spin-locks

- CPU Profiling tool
  - Takes a snapshot of the CPU every millisecond
    - Interval is a parameter
  - At each snapshot, it records what function is running
  - After the process has finished, we can look at the log
  - If 100 snapshots were taken in function X, we can estimate that function X ran for 100ms
- What do we measure for spin-locks?

# CPU profiling tool - gperftools

- Installation

- <https://github.com/gperftools/gperftools/releases>
- Unpack and cd to directory
- `./configure --prefix=<path>`
- `make`
- `make install`
- `(make clean)`

The screenshot shows the GitHub release page for gperftools-2.8. On the left, there's a sidebar with a 'Latest release' badge, the repository name 'gperftools-2.8', a commit hash '180bfa1', and a 'Compare' dropdown. The main content area has the title 'gperftools-2.8' and a message from 'alk' stating 'gperftools 2.8 is out!' and 'Here are notable changes:'. The changes list includes: 'ProfilerGetStackTrace is now officially supported', 'Build failures on mingw were fixed. This fix is included in the 2.8 release', 'Build failure of page\_heap\_test on MSVC is fixed', 'Ryan Macnak contributed fix for compiling on Windows', and 'test failures caused by new gcc 10 optimizations are fixed'. Below the changes, there's an 'Assets' section with 4 items. The first item, 'gperftools-2.8.tar.gz', is highlighted with a red box and a red arrow pointing to it from the left. Other assets include 'gperftools-2.8.zip', 'Source code (zip)', and 'Source code (tar.gz)'.

Latest release

gperftools-2.8

180bfa1

Compare

## gperftools-2.8

alk released this on Jul 6 · 2 commits to master si

gperftools 2.8 is out!

Here are notable changes:

- ProfilerGetStackTrace is now officially supported
- Build failures on mingw were fixed. This fix is included in the 2.8 release
- Build failure of page\_heap\_test on MSVC is fixed
- Ryan Macnak contributed fix for compiling on Windows
- test failures caused by new gcc 10 optimizations are fixed

Assets 4

- gperftools-2.8.tar.gz
- gperftools-2.8.zip
- Source code (zip)
- Source code (tar.gz)

# CPU profiling tool - gperftools

- Usage
  - Run your executable with the CPUPROFILE environment variable set
    - LD\_PRELOAD=<path/to/libprofiler> CPUPROFILE=<path/to/profile\_output>  
<path/to/binary> [binary args]
  - Run pprof to analyze the CPU usage
    - pprof <path/to/binary> <path/to/profile\_output>
    - Profiling report will be printed to stdout

# CPU profiling tool - gperftools

- Usage

- Run your executable with the LD\_PRELOAD and CPUPROFILE environment variable set
  - LD\_PRELOAD=<path/to/libprofiler>  
CPUPROFILE=<path/to/profile\_output>  
t> <path/to/binary> [binary args]
- Run pprof to analyze the CPU usage
  - pprof <path/to/binary>  
<path/to/profile\_output>
  - Profiling report will be printed to stdout

- Example

- Run your executable with the LD\_PRELOAD and CPUPROFILE environment variable set
  - LD\_PRELOAD=/usr/lib/libprofiler.so  
CPUPROFILE=raw.perf lab2\_list --  
threads=12 --iterations=1000 --  
sync=s
- Run pprof to analyze the CPU usage
  - pprof lab2\_list raw.perf

# CPU profiling tool - gperftools

- Output
  - Obtained from
    - `pprof --text <path_to_binary> <path_to_profile_output>`
  - Text output will produce lines that look like
  - `14 2.1% 17.2% 58 8.7% std::_Rb_tree::find`

# CPU profiling tool - gperftools

- Output
  - Obtained from
    - `pprof --text <path_to_binary> <path_to_profile_output>`
  - Text output will produce lines that look like
  - `14 2.1% 17.2% 58 8.7% std::_Rb_tree::find`
  - # samples (ms) spent in this function
  - % of samples in this function
  - Cumulative % of samples spent so far
    - % on this line + % on all previous lines. Last line shows 100%.
  - # samples in this function + callees
  - % of samples spent in this function + its callees
  - Function name

# CPU profiling tool - gperftools

- Sampling result
  - Obtained from
    - `pprof --list=<func_name> <path_to_binary> <path_to_profile_output>`

Sample  
count

1622	1622	17:	for (i = 0; i < MAX_SUM/thread_num; i++)
.	.	18:	{
13794	13794	19:	while (__sync_lock_test_and_set(&lock, 1));
769	769	20:	sum++;
2	2	21:	__sync_lock_release(&lock);
.	.	22:	}



# Measuring the overhead for mutexes

- Can we use gperftools to measure?

# Measuring the overhead for mutexes

- We can't use gperftools! Why?
  - When a mutex can't be acquired, the thread goes to sleep.
  - The time spent sleeping can't be measured by counting CPU cycles.
- How do we measure the overhead for mutexes?

# Measuring the overhead for mutexes

- We can't use gperftools! Why?
  - When a mutex can't be acquired, the thread goes to sleep.
  - The time spent sleeping can't be measured by counting CPU cycles.
- How do we measure the overhead for mutexes?
  - Directly measure the time taken to acquire a mutex

# Timing mutex waits

- Computing wait-for-lock time
  - Measure time before and after getting the lock
  - Add up all wait time for all threads
  - Divide by number of operations
  - Output statistics for the run

# Discovering and addressing the problem

- The timings allow you to prove the original intuition: degradation is a result of increased contention
- To decrease contention, we can split the list
  - Add a `--lists` option to specify the number of sub-lists
- Each thread does:
  - Insert *#iterations* nodes to the multi-list
    - Which sub-list to insert in is determined by a hash function
  - Look-up and delete inserted nodes

# Inserting into different sub-lists

- No requirement of total ordering
  - Each sub-list must be ordered
  - But there is no order requirement between sub-lists
- Selecting sub-list
  - Use a hash function to map value  $\rightarrow$  sub-list
  - Doesn't matter which hash function you use
  - Could be as simple as a modulo operation

Why is this more efficient?

# Why is this more efficient?

- Splitting lists increases efficiency by decreasing contention
  - Not all threads try to acquire the same lock
- What else?



# Why is this more efficient?

- Splitting lists increases efficiency by decreasing contention
  - Not all threads try to acquire the same lock
- The sub-lists are shorter!
  - Insertion operations ( $O(N)$ ) will be faster
  - The time spent in the critical section will decrease
  - All threads, on average, wait a shorter time to acquire a lock
  - Fewer cycles wasted to acquire a lock

# In this project you will...

- Modify lab2a to support
  - insertion and deletion to multiple sub-lists
- Time synchronization cost
  - Using profiling tool and regular timing function
- Evaluate profiling report and answer questions

Questions?