

# CS 111 week 5

## Project 2b: Lock contention

Discussion 1B

Tianxiang Li

# Lock contention: the problem

```
#define MAX_SUM 100000000
```

```
void * thread_worker(void * unused)
```

```
{
```

```
    int i = 0;
```

```
    for (i = 0; i < MAX_SUM/thread_num; i++) //thread_num is the total number of threads
```

```
    {
```

```
        while (___sync_lock_test_and_set(&lock, 1));
```

```
        sum++;
```

```
        ___sync_lock_release(&lock);
```

```
    }
```

```
}
```

Will we gain speedup if we increase the number of threads?

*Parallel threads are not able to run in parallel if they concurrently attempt to access a shared lock.*

# Project 2B: Overview

## Part A: Identify the lock contention problem

1. Starting from lab2a\_list, **measure throughput** (# of total ops/seconds)
2. Use google profile tool to **pinpoint** which the line causing the **problem** (spin-lock)
3. **Measure the time** spent on acquiring mutex (mutex-lock)

## Part B: Resolve the lock contention problem:

Implement the linked list with **N sublists**.

--enable **parallel access to the list**

# Submission Tarball

- **sortedList.h** (do not modify)
  - linked list operation interfaces, must implement
- **sortedList.c**
  - implements insert, delete, lookup, and length methods for a sorted doubly linked list
- **lab2\_list.c**
  - --threads, --iterations, --yield, --sync, --lists
  - Parallel threads operate on linked list, output performance result
- **profile.out**
  - profiling report: where time was spent in the un-partitioned spin-lock implementation
- **lab2b\_list.csv**
- **graphs** (5 png)
- **README** (included files, refs, other info, e.g. research, limitations, features, testing methodology)
- **Scripts** (lab2b.gp)
- **Makefile**
  - **default** ... compile executable ( **-Wall -Wextra** options)
  - **tests** ... run tests cases, generate CSV results
  - **profile** ... run tests with profiling tools (google-pprof)
  - **graphs** ... generate graphs (gnuplot lab2b.gp)
  - **dist** ... \*.c,\*.h,\*.gp, \*.csv,\*.out,\*.png, README, Makefile
  - **clean** ... tarball,exec,\*.o

# Review Your results from previous Lab

- All the locks experience exponential rise of **cost/operation** vs **#threads**
- Reason
  - *Spin lock is using cycles*
  - *Mutex lock have increased wait time*

# Testing

- In a single-threaded implementation, time per operation is a very reasonable performance metric.
- But for **multi-threaded** implementations we are also concerned about how well we are **taking advantage of the available parallelism**.
- This is more clearly seen in the **aggregate throughput** (total operations per second for all threads combined).
  - **Mutex** synchronized list operations, 1,000 iterations, 1,2,4,8,12,16,24 threads
  - **Spin-lock** synchronized list operations, 1,000 iterations, 1,2,4,8,12,16,24 threads
  - Capture the output, and produce a plot of **the total number of operations per second** for each synchronization method

# Part A: Using the Profile Tool

# Profile Tool: motivation

```
#define MAX_SUM 100000000
```

```
void * thread_worker(void * unused)
```

```
{
```

```
    int i = 0;
```

```
    for (i = 0; i < MAX_SUM/thread_num; i++) //thread_num is the total number of threads
```

```
    {
```

```
        while (__sync_lock_test_and_set(&lock, 1));
```

```
        sum++;
```

```
        __sync_lock_release(&lock);
```

```
    }
```

```
}
```

How can we design a tool to help us find the bottleneck of the program?

(which function/LOC takes the most time)

Ideally such a tool needs to be

1. Transparent: Does not require modifications to the program
2. Low overhead
3. Easy to enable/disable



# Profile Tool Key idea: Sampling

Key idea of profile tool: [sampling](#)

For every N (say 1,000,000) instructions/cycles, check the program counter, see which function/LOC it falls in. Aggregate the results when the program exits.

Sampling tool satisfies all the three goals we have mentioned:

- Transparent,

- Low-overhead: Sampling rate is low

- Easy to enable/disable: Explained later

# Execution Profiling

**Analyze a program's execution to determine how much time is being spent in which routines**

- The standard Linux *gprof(1)* tool
  - is quite simple to use, but its **call-counting mechanism** is **not-multi-thread safe**, and its execution sampling is not multi-thread aware.
  - not usable for analyzing the performance of multi-threaded applications.
- *valgrind*
  - best known for its memory leak detector, which has an interpreted execution engine that can extract a great deal of information about where CPU time is being spent, even estimating cache misses.
  - It does work for multi-threaded programs, but its **interpreter does not provide much parallelism**. As such it is not useful for examining high contention situations.
- *gperftools*
  - a wonderful set of performance optimization tools from Google. It includes a profiler that is quite similar to *gprof* (in that it samples real execution). This is probably the **best tool to use for this problem**.

# Google Profile Tool: Installation

## Using package management tool:

With root permission:

*apt-get install google-perftools* (Ubuntu)

*yum install gperftools* (Old Version of CentOS, Fedora)

*dnf install gperftools* (New Version of CentOS, Fedora)

## Install from source code:

*git clone <https://github.com/gperftools/gperftools>*

*./autogen.sh*

*./configure* # use --prefix option to set installation path if needed

*make*

*make install*

*Note: First check whether Inxsrv09 already has installed gperftools*

# Google Profile Tool: pprof

- pprof **reads** a collection of **profiling samples** and **generates reports** to visualize data
  - run the spin-lock list test (1,000 iterations 12 threads) under the profiler
- Usage
  - **google-pprof** [options] <program> <profile>
    - options
      - **--text** Generate text report [default]: Shows you which routine is consuming most of the CPU time
      - **--list=<routine>** Generate source listing of matching routines: Give you a source-level breakdown of how much time is being spent on each instruction
    - program
      - ./lab2\_list
    - profile
      - ./raw.gperf

# Google Profile Tool: Usage

1. **Compile** your program with **-g**
2. Find where you have installed **libprofiler.so** and **pprof**, the below example assume `~/lib/libprofiler.so` (profiling library) and `~/bin/pprof` (user level app)
3. **LD\_PRELOAD**=`~/lib/libprofiler.so`  
**CPUPROFILE**=`./raw.gperf ./lab2_list --threads=12 --iterations=1000 --sync=s`  
*# LD\_PRELOAD: link the library, overload existing function in my\_prog*  
*# CPUPROFILE: gperf specific environment variable, tells gperf where to store the raw profiling data. (perform sampling on ./lab2\_list output to ./raw.gperf)*
4. `~/bin/pprof --text ./lab2_list ./raw.gperf`  
*#analyze raw sampling file, show sampling results in functions (i.e. how often does sampling fall into individual functions) step 4,5 generates readable output from raw output*
5. `~/bin/pprof --list=thread_worker ./lab2_list ./raw.gperf`  
*# show sampling results in function thread\_worker (i.e. how often does sampling fall into each lines of code in thread\_worker)*

# Analyzing text output

`~/bin/pprof --text ./myprog ./raw.gperf`

Total: 16187 samples

Text mode has lines of output that look like this:

```
16187  100.0%  100.0%  16187  100.0%.  thread_worker
```

Here is how to interpret the columns:

- Number of profiling samples this function was running in
- Percentage of profiling samples this function was running in
- Percentage of profiling samples in the functions printed so far
- Number of profiling samples in this function and its callees
  - The number of samples in which the function appeared (either running or waiting for a called function to return)
- Percentage of profiling samples in this function and its callees
- Function name

# Sampling results

*~/gperf/bin/pprof --list=thread\_worker ./test ./raw.gperf*

Sample  
count

1622	1622	17:	for (i = 0; i < MAX_SUM/thread_num; i++)
.	.	18:	{
13794	13794	19:	while (__sync_lock_test_and_set(&lock, 1));
769	769	20:	sum++;
2	2	21:	__sync_lock_release(&lock);
.	.	22:	}

# Timing Mutex Waits

- **spin-lock**

- profiling can tell us **where we are spending most of our time.**

- **mutex**

- a thread that cannot get the lock is **blocked, does not consume CPU time.**
  - Profiling only tells us **what code we are executing.** It doesn't tell us anything about the **time we spend** blocked.
  - How could we confirm that, in the mutex case, most threads are spending most of their time waiting for a lock?



## Part A: Measuring the lock contention time

# Code snippet of thread\_worker from lab2a

```
void * thread_worker(threadNum) {  
    startIndex = threadNum * iteration;  
    for (i = startIndex; i < startIndex + iteration; i++)  
    {  
        pthread_mutex_lock(&list_lock);  
        SortedList_insert(listhead, pool[i]);  
        pthread_mutex_unlock(&list_lock);  
    }  
    ....  
}
```

# Measure lock contention time

```
void * thread_worker(threadNum) {  
    startIndex = threadNum * iteration;  
    for (i = startIndex; i < startIndex + iteration; i++)  
    {  
        clock_gettime(CLOCK_MONOTONIC, &start_time);  
        pthread_mutex_lock(&list_lock);  
        clock_gettime(CLOCK_MONOTONIC, &end_time);  
  
        SortedList_insert(listhead, pool[i]);  
        pthread_mutex_unlock(&list_lock);  
  
        //wait_time is an unsigned long array with length being number of worker threads.  
        wait_time[threadNum] += calc_diff(&start_time, &end_time);  
    }  
    ....  
}
```

# Addressing the Underlying Problem

- Fundamental problem
  - Throughput degrade is the result of **increased contention**.
- Classic solution
  - **partition the single resource** (in this case a linked list) into multiple independent resources
  - divide the requests among those sub-resources

Part B: Break linked lists into sublists

# Part B requirement

- Add **--lists=#** option
  - **break the single sorted list** into the specified number of sub-lists
  - each with its own list header and synchronization object
- select which sub-list a particular key should be in
  - based on a simple hash of the key, modulo the number of lists
- obtain the length of the list
  - enumerating all of the sub-lists.
- each thread:
  - starts with a set of pre-allocated and **initialized elements** (--iterations=#)
  - **inserts** them all into the multi-list (which sublist the key should go into determined by a hash of the key)
  - gets the **list length**
  - **looks up** and **deletes** each of the keys it inserted
  - exits to re-join the parent thread

## Part B: overview

```
SortedListElement_t listheads[], * pool;
pthread_mutex_t locks[];
void * thread_worker(threadNum) {
    startIndex = threadNum * iteration;
    for (i = startIndex; i < startIndex + iteration; i++) Mul_SortedList_insert(pool[i]);
    Mul_SortedList_length();
    for (i = startIndex; i < startIndex + iteration; i++) {
        e = Mul_SortedList_lookup(pool[i]->key);
        SortedList_delete(e);
    }
}
```

# Insert of multiple sublists

Note: below is pseudocode

SortedListElement\_t listheads[]; //one listhead for each sublist

pthread\_mutex\_t locks[]; //one lock for each sublist

```
void Mul_SortedList_insert(SortedListElement_t *element)
{
    int list_num = hashkey(element->key); //hashkey return between [0 , # of sublists-1]
    pthread_mutex_lock(&locks[list_num]);
    SortedList_Insert(&listheads[list_num], element);
    pthread_mutex_unlock(&locks[list_num]);
}
```



# Lookup of multiple sublists

Note: below is pseudocode

```
SortedListElement_t listheads[];
```

```
pthread_mutex_t locks[];
```

```
SortedListElement_t * Mul_SortedList_lookup(char * key)
```

```
{
```

```
    SortedListElement_t * ret = NULL;
```

```
    int list_num = hashkey(key); //determine which sublist the element of key belongs to
```

```
    pthread_mutex_lock(&locks[list_num]);
```

```
    ret = SortedList_lookup(&listheads[list_num], key); //get the element of the key
```

```
    pthread_mutex_unlock(&locks[list_num]);
```

```
    return ret;
```

```
}
```

# Length of multiple sublists: naïve implementation

Note: below is pseudocode

```
SortedListElement_t listheads[];
```

```
pthread_mutex_t locks[];
```

```
int Mul_SortedList_length(void)
```

```
{
```

```
    int i = 0, length = 0;
```

```
    for (i = 0; i < num_sub_lists; i++)
```

```
    {
```

```
        pthread_mutex_lock(&locks[i]);
```

```
        length += SortedList_length(listheads[i]);
```

```
        pthread_mutex_unlock(&locks[i]);
```

```
    }
```

```
    return length;
```

```
}
```

Drawing Figures with gnuplot

```
set terminal png
```

```
set datafile separator ", "
```

```
set title "List-1: Cost per Operation vs Iterations"
```

```
set xlabel "Iterations"
```


```
set logscale x 10
```

```
set ylabel "Cost per Operation (ns)"
```

```
set logscale y 10
```

```
set output 'lab2_list-1.png'
```

```
plot \
  "< grep 'list-none-none,1,' lab2_list.csv" \
    using ($3):($7) title 'raw' with \
    linespoints lc rgb 'red', \
  "< grep 'list-none-none,1,' lab2_list.csv" \
    using ($3):($7)/(4*($3)) \
    title '/4 x iterations' with linespoints lc rgb 'green'
```



```
list-none-none,1,10,1,30,139945,4664
```

```
list-none-none,1,100,1,300,165364,551
```

```
list-none-none,1,1000,1,3000,3928455,1309
```

```
list-none-none,1,10000,1,30000,495197933,16506
```

```
list-none-none,1,20000,1,60000,3054392526,50906
```

```
list-none-none,1,1000,1,3000,3879764,1293
```